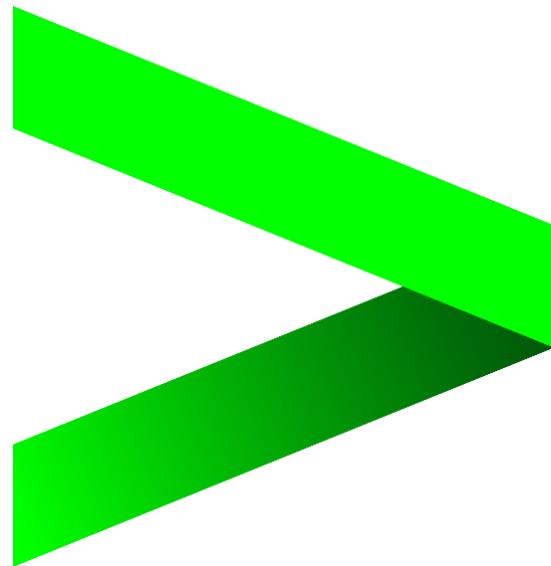


## Hands-on Workbook

### Kubernetes



Version	Revision date	Description	Author	Reviewed by	Approved by
1.0	1-Jun-2018	Initial version	Srinivas.c.mamidi	sreenivasa.cpcl.rao	sreenivasa.cpcl.rao

---

*Table of Contents*

---

1.	Exercise 1.1 Setup Docker development environment on a Linux VM .....	3
2.	Exercise 1.2 Building an app in the Docker way .....	5
3.	Exercise 2.1 Setting up single node Kubernetes cluster using Minikube.....	12
4.	Exercise 3.1 Deploying an App to Kubernetes cluster .....	15
5.	Exercise 3.2 Exploring App Deployed in Kubernetes cluster .....	17
6.	Exercise 3.3 Exposing deployed app using services.....	20
7.	Exercise 3.4 Scaling your App .....	22
8.	Exercise 3.5 Updating App .....	27
9.	Exercise 4.1 Deploying FirstHello app onto Kubernetes cluster .....	33
10.	Exercise 4.2 Deploying webserver on to Kubernetes cluster .....	36
11.	Exercise 4.3 Deploying multi-tier app on to Kubernetes cluster .....	45
12.	Exercise 4.4 Deploying MySQL onto Kubernetes cluster with persistent volumes (optional) .....	60
13.	Exercise x.x Setup Kubernetes cluster using Kubeadm .....	64

---

*Module 1 – Docker setup and basics*

---

## Exercise 1.1 Setup Docker development environment on a Linux VM

### Scenario

Docker setup on Windows 10 machine

### To do tasks:

In this exercise

- Docker setup on Linux
- Testing Docker setup on Linux

### Steps

Docker setup on windows	
1	<p>Use below steps to <b>install Docker</b> on Linux (CentOS) environment</p> <p>Download Docker latest CentOS version from “<a href="https://download.docker.com/linux/centos/7/x86_64/stable/Packages/">https://download.docker.com/linux/centos/7/x86_64/stable/Packages/</a>”</p> <div style="border: 1px solid #ccc; padding: 5px;"><pre>File Edit View Search Terminal Help [user1@ip-172-42-32-4 Downloads]\$ ls <b>docker-ce-18.03.1.ce-1.el7.centos.x86_64.rpm</b> [user1@ip-172-42-32-4 Downloads]\$</pre></div>
2	<p>Go to directory “/home/user1/Downloads”</p> <p>Run command “<code>yum install docker-ce-18.03.1.ce-1.el7.centos.x86_64.rpm</code>” as shown below to install Docker</p>

## Container orchestration with Kubernetes Hands-on Workbook

	<pre>[user1@ip-172-42-0-39 Downloads]\$ sudo yum install docker-ce-18.03.1.ce-1.el7.centos.x86_64.rpm Loaded plugins: fastestmirror, langpacks Examining docker-ce-18.03.1.ce-1.el7.centos.x86_64.rpm: docker-ce-18.03.1.ce-1.el7.centos.x86_64 Marking docker-ce-18.03.1.ce-1.el7.centos.x86_64.rpm to be installed Resolving Dependencies --&gt; Running transaction check --&gt; Package docker-ce.x86_64 0:18.03.1.ce-1.el7.centos will be installed --&gt; Processing Dependency: container-selinux &gt;= 2.9 for package: docker-ce-18.03.1.ce-1.el7.centos.x86_64 Loading mirror speeds from cached hostfile  * base: mirrors.fibergrid.in  * epel: del-repos.extreme-ix.org  * extras: mirrors.fibergrid.in  * nux-dextop: mirror.li.nux.ro  * updates: mirrors.fibergrid.in --&gt; Processing Dependency: pigz for package: docker-ce-18.03.1.ce-1.el7.centos.x86_64 --&gt; Running transaction check --&gt; Package container-selinux.noarch 2:2.55-1.el7 will be installed --&gt; Processing Dependency: selinux-policy-targeted &gt;= 3.13.1-183 for package: 2:container-selinux-2.55-1.el7.noarch --&gt; Processing Dependency: selinux-policy-base &gt;= 3.13.1-183 for package: 2:container-selinux-2.55-1.el7.noarch --&gt; Processing Dependency: selinux-policy &gt;= 3.13.1-183 for package: 2:container-selinux-2.55-1.el7.noarch --&gt; Package pigz.x86_64 0:2.3.4-1.el7 will be installed --&gt; Running transaction check --&gt; Package selinux-policy.noarch 0:3.13.1-166.el7_4.7 will be updated --&gt; Package selinux-policy.noarch 0:3.13.1-192.el7_5.4 will be an update --&gt; Processing Dependency: policycoreutils &gt;= 2.5-18 for package: selinux-policy-3.13.1-192.el7_5.4.noarch --&gt; Package selinux-policy-targeted.noarch 0:3.13.1-166.el7_4.7 will be updated --&gt; Package selinux-policy-targeted.noarch 0:3.13.1-192.el7_5.4 will be an update --&gt; Running transaction check --&gt; Package policycoreutils.x86_64 0:2.5-17.1.el7 will be updated --&gt; Processing Dependency: policycoreutils = 2.5-17.1.el7 for package: policycoreutils-python-2.5-17.1.el7.x86_64 --&gt; Package policycoreutils.x86_64 0:2.5-22.el7 will be an update Verifying : policycoreutils-python-2.5-22.el7.x86_64 9/25 Verifying : 2:container-selinux-2.55-1.el7.noarch 10/25 Verifying : libsepolicy-2.5-8.1.el7.x86_64 11/25 Verifying : docker-ce-18.03.1.ce-1.el7.centos.x86_64 12/25 Verifying : libselinux-2.5-12.el7.x86_64 13/25 Verifying : libselinux-utils-2.5-12.el7.x86_64 14/25 Verifying : libselinux-utils-2.5-11.el7.x86_64 15/25 Verifying : libselinux-2.5-11.el7.x86_64 16/25 Verifying : libsepolicy-2.5-6.el7.x86_64 17/25 Verifying : policycoreutils-python-2.5-17.1.el7.x86_64 18/25 Verifying : policycoreutils-2.5-17.1.el7.x86_64 19/25 Verifying : libsemanage-python-2.5-8.el7.x86_64 20/25 Verifying : selinux-policy-3.13.1-166.el7_4.7.noarch 21/25 Verifying : selinux-policy-targeted-3.13.1-166.el7_4.7.noarch 22/25 Verifying : libsemanage-2.5-8.el7.x86_64 23/25 Verifying : libselinux-python-2.5-11.el7.x86_64 24/25 Verifying : setools-libs-3.3.8-1.1.el7.x86_64 25/25  Installed:   docker-ce.x86_64 0:18.03.1.ce-1.el7.centos  Dependency Installed:   container-selinux.noarch 2:2.55-1.el7          pigz.x86_64 0:2.3.4-1.el7  Dependency Updated:   libselinux.x86_64 0:2.5-12.el7                libselinux-python.x86_64 0:2.5-12.el7    libselinux-utils.x86_64 0:2.5-12.el7   libsemanage.x86_64 0:2.5-11.el7                libsemanage-python.x86_64 0:2.5-11.el7  libsepolicy.x86_64 0:2.5-8.1.el7   policycoreutils.x86_64 0:2.5-22.el7           policycoreutils-python.x86_64 0:2.5-22.el7  selinux-policy.noarch 0:3.13.1-192.el7_5.4   selinux-policy-targeted.noarch 0:3.13.1-192.el7_5.4  setools-libs.x86_64 0:3.3.8-2.el7  Complete! [user1@ip-172-42-0-39 Downloads]\$</pre>
3	Run command “ <code>systemctl start docker</code> ” to start Docker
4	Run command “ <code>docker run hello-world</code> ” to see everything working

## Container orchestration with Kubernetes Hands-on Workbook

```
[user1@ip-172-42-0-39 ~]$ sudo docker run hello-world
Unable to find image 'hello-world:latest' locally
latest: Pulling from library/hello-world
9db2ca6ccae0: Pull complete
Digest: sha256:4b8ff392a12ed9ea17784bd3c9a8b1fa3299cac44aca35a85c90c5e3c7afacdc
Status: Downloaded newer image for hello-world:latest
```

Hello from Docker!

This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:

1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.  
(amd64)
3. The Docker daemon created a new container from that image which runs the executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it to your terminal.

To try something more ambitious, you can run an Ubuntu container with:

```
$ docker run -it ubuntu bash
```

Share images, automate workflows, and more with a free Docker ID:

<https://hub.docker.com/>

For more examples and ideas, visit:

<https://docs.docker.com/engine/userguide/>

```
[user1@ip-172-42-0-39 ~]$ █
```

Run command "docker images" as shown below

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
hello-world	latest	fce289e99eb9	4 weeks ago	1.84kB

## Exercise 1.2 Building an app in the Docker way

### Scenario

Building a web app in the Docker way

### To do tasks:

In this exercise, you will be:

- Writing dockerfile to containerize app, image creation
- Running containers, tagging and pushing to central repository's
- Pulling docker images from central repository

### Steps

Containerizing a Python app	
1	Create new workspace as shown below  [user1@ip-172-42-32-7 DockerWS]\$ pwd /kubews/DockerWS [user1@ip-172-42-32-7 DockerWS]\$ █
2	Understand Docker file as shown below

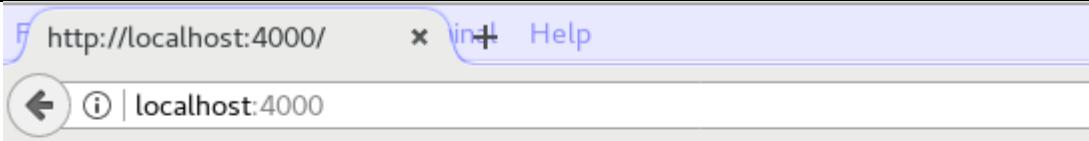
	<pre>[root@ip-172-42-32-4 DockerWS]# ls app.py  Dockerfile  requirements.txt [root@ip-172-42-32-4 DockerWS]# cat Dockerfile # Use an official Python runtime as a parent image  FROM python:2.7-slim  # Set the working directory to /app  WORKDIR /app  # Copy the current directory contents into the container at /app  ADD . /app  # Install any needed packages specified in requirements.txt  RUN pip install --trusted-host pypi.python.org -r requirements.txt  # Make port 80 available to the world outside this container  EXPOSE 80  # Define environment variable Acc LKM  ENV NAME World  # Run app.py when the container launches  CMD ["python", "app.py"]  # https://docs.docker.com/engine/reference/builder/#dockerignore-file [root@ip-172-42-32-4 DockerWS]# █</pre>
3	Understand app.py file with below mentioned content

## Container orchestration with Kubernetes Hands-on Workbook

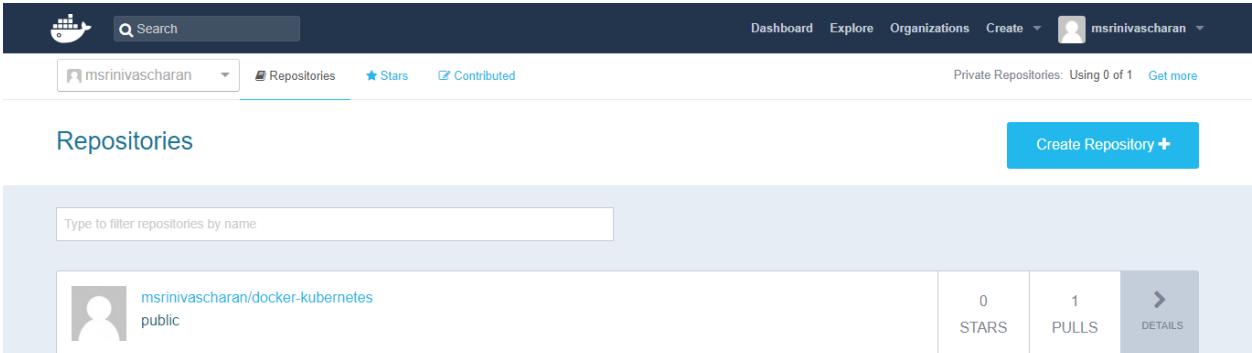
	<pre>[root@ip-172-42-32-4 DockerWS]# cat app.py from flask import Flask from redis import Redis, RedisError import os import socket  # Connect to Redis redis = Redis(host="redis", db=0, socket_connect_timeout=2, socket_timeout=2)  app = Flask(__name__)  @app.route("/") def hello():     try:         visits = redis.incr("counter")     except RedisError:         visits = "&lt;i&gt;cannot connect to Redis, counter disabled&lt;/i&gt;"      html = "&lt;h3&gt;Hello LKM Accenture {name}!&lt;/h3&gt; \         "&lt;b&gt;Hostname:&lt;/b&gt; {hostname}&lt;br/&gt;" \         "&lt;b&gt;Visits:&lt;/b&gt; {visits}"     return html.format(name=os.getenv("NAME", "world-Accenture LKM 28th Sep18"), hostname=socket.gethostname(), visits=visits)  if __name__ == "__main__":     app.run(host='0.0.0.0', port=80) [root@ip-172-42-32-4 DockerWS]#</pre>
4	<p>Understand requirements.txt with below mentioned content</p> <pre>[root@ip-172-42-32-4 DockerWS]# cat requirements.txt Flask  Redis [root@ip-172-42-32-4 DockerWS]#</pre>
5	<p>Top level of your new directory must contain files as shown below</p> <pre>[user1@ip-172-42-32-7 DockerWS]\$ ls app.py  Dockerfile  requirements.txt [user1@ip-172-42-32-7 DockerWS]\$ █</pre>
6	<p>Run the build command “<code>docker build -t firsthello .</code>” from “<code>/kubews/DockerWS</code>” to creates a Docker image as shown below</p>

## Container orchestration with Kubernetes Hands-on Workbook

	<pre> Step 1/7 : FROM python:2.7-slim 2.7-slim: Pulling from library/python 5e6ec7f28fb7: Pull complete 76e195ed530a: Pull complete 674dec4e4375b: Pull complete c306b9f6d5e3: Pull complete Digest: sha256:9be0e36b98223ee70e92a610991e3684416a70ae6b5afe22f818b675ee22e014 Status: Downloaded newer image for python:2.7-slim --&gt; 4620f1f365b9 Step 2/7 : WORKDIR /app Removing intermediate container 0cb18892f88b --&gt; 8061a35caa9e Step 3/7 : ADD . /app --&gt; e4ac4fbdaaff3 Step 4/7 : RUN pip install --trusted-host pypi.python.org -r requirements.txt --&gt; Running in fca4c8374d20 DEPRECATION: Python 2.7 will reach the end of its life on January 1st, 2020. Please upgrade your Python as Python 2.7 won't be maintained after that date. A future version Collecting Flask (from -r requirements.txt (line 1)) --&gt; Downloading https://files.pythonhosted.org/packages/7f/e7/0857774ed4536d3242b14dacb4696386634607af824ea997202cd0edb4b/Flask-1.0.2-py2.py3-none-any.whl (91kB) Collecting Redis (from -r requirements.txt (line 3)) --&gt; Downloading https://files.pythonhosted.org/packages/f1/19/a028b77c23f9fdbcc6480787a60807c78a45947593a02dbf026636c90d/redis-3.1.0-py2.py3-none-any.whl (63kB) Collecting itsdangerous==0.24 (from Flask-&gt;-r requirements.txt (line 1)) --&gt; Downloading https://files.pythonhosted.org/packages/76/ae/44b03b253d6fade317f32c24d10b3b35c2239807046a4c953c7b89fa49e/itsdangerous-1.1.0-py2.py3-none-any.whl Collecting Jinja2==2.10 (from Flask-&gt;-r requirements.txt (line 1)) --&gt; Downloading https://files.pythonhosted.org/packages/7f/ff/ae64bacdfc95f27a016a7bed8e686763ba4d277a76f326592208731/Jinja2-2.10-py2.py3-none-any.whl (126kB) Collecting Werkzeug==0.14 (from Flask-&gt;-r requirements.txt (line 1)) --&gt; Downloading https://files.pythonhosted.org/packages/20/c4/12e3e56473e52375aa29c4764e70d1b8f3efa6682bef8d0aae04fe335243/Werkzeug-0.14.1-py2.py3-none-any.whl (322kB) Collecting click==5.1 (from Flask-&gt;-r requirements.txt (line 1)) --&gt; Downloading https://files.pythonhosted.org/packages/fa/37/45185cb5abb30d7257104c434fe0b07e5a195a6847506c074527aa599ec/Click-7.0-py2.py3-none-any.whl (81kB) Collecting MarkupSafe==0.23 (from Jinja2==2.10-&gt;Flask-&gt;-r requirements.txt (line 1)) --&gt; Downloading https://files.pythonhosted.org/packages/bc/3a/6bffd7b4b202f7a33bdda8e4e3d3acc719f381fd730f9a0e7c5f34e845bd4d/MarkupSafe-1.1.0-cp27-cp27mu-manylinux1_x86_64.whl Installing collected packages: itsdangerous, MarkupSafe, Jinja2, Werkzeug, click, Flask, Redis Successfully installed Flask-1.0.2 Jinja2-2.10 MarkupSafe-1.1.0 Redis-3.1.0 Werkzeug-0.14.1 click-7.0 itsdangerous-1.1.0 Removing intermediate container fca4c8374d20 --&gt; 32c8b73a989c Step 5/7 : EXPOSE 80 --&gt; Running in cf343cf6b1a Removing intermediate container cf343cf6b1a --&gt; 2656778db116 Step 6/7 : ENV NAME World --&gt; Running in 029e282004a7 Removing intermediate container 029e282004a7 --&gt; 1509300165a9 Step 7/7 : CMD ["python", "app.py"] --&gt; Running in 41528f284346 Removing intermediate container 41528f284346 --&gt; 0b25d2c5ab0d Successfully built 0b25d2c5ab0d Successfully tagged firsthello:latest [root@ip-172-42-32-4 DockerWS]# </pre>
7	<p>Check local Docker image registry using “<code>docker image ls</code>” to see newly created app image.</p> <pre> [root@ip-172-42-32-4 DockerWS]# docker image ls REPOSITORY          TAG           IMAGE ID            CREATED             SIZE firsthello          latest        0b25d2c5ab0d    2 minutes ago      131MB python              2.7-slim     4620f1f365b9    6 days ago         120MB hello-world         latest        fce289e99eb9    4 weeks ago        1.84kB [root@ip-172-42-32-4 DockerWS]# </pre>
8	<p>Run the app, mapping your machine’s port 4000 (or any free port) to the container’s published port 80 using -p. command is “<code>docker run -p 4000:80 firsthello</code>”</p> <pre> [user1@ip-172-42-32-7 DockerWS]\$ sudo docker run -p 4000:80 firsthello  * Serving Flask app "app" (lazy loading)  * Environment: production    WARNING: Do not use the development server in a production environment.    Use a production WSGI server instead.  * Debug mode: off  * Running on http://0.0.0.0:80/ (Press CTRL+C to quit) 172.17.0.1 - - [31/Jul/2018 07:02:23] "GET / HTTP/1.1" 200 - 172.17.0.1 - - [31/Jul/2018 07:02:23] "GET /favicon.ico HTTP/1.1" 404 - 172.17.0.1 - - [31/Jul/2018 07:04:08] "GET / HTTP/1.1" 200 - </pre> <p>Go to web client, type URL <a href="http://localhost:4000">http://localhost:4000</a> to see app UI</p>

	 <p><b>Hello LKM Accenture World!</b></p> <p><b>Hostname:</b> 69153b67e897  <b>Visits:</b> <i>cannot connect to Redis, counter disabled</i></p>
9	<p>You can run app in the background, in detached mode as shown below.</p> <pre>[user1@ip-172-42-32-7 DockerWS]\$ sudo docker run -d -p 4000:80 firsthello c62e5a65f534f299cd7a4bd0fb3e7bc7de1317aaad1be6cef451a68e6fa9264b [user1@ip-172-42-32-7 DockerWS]\$</pre> <p>You can run command “<code>docker container stop &lt;container id&gt;</code>” to stop container as shown below</p> <pre>[user1@ip-172-42-1-80 ~]\$ sudo docker ps CONTAINER ID        IMAGE               COMMAND                  CREATED             STATUS              PORTS                 NAMES e035b0956e6e        firsthello          "python app.py"         8 minutes ago      Up 8 minutes       0.0.0.0:4000-&gt;80/tcp   agitated_kapitsa [e035b0956e6e] [user1@ip-172-42-1-80 ~]\$ sudo docker stop e035b0956e6e [user1@ip-172-42-1-80 ~]\$ sudo docker ps CONTAINER ID        IMAGE               COMMAND                  CREATED             STATUS              PORTS                 NAMES [e035b0956e6e]</pre>
10	<p>Note: Before this step ensure that you have access to <a href="https://hub.docker.com/">https://hub.docker.com/</a></p> <p>Log in to the Docker public registry on your local machine as shown below</p> <pre>[root@ip-172-42-32-4 DockerWS]# docker login Login with your Docker ID to push and pull images from Docker Hub. If you don't have a Docker ID, head over to https://hub.docker.com to create one. Username: msrinivascharan Password: Login Succeeded [root@ip-172-42-32-4 DockerWS]#</pre> <p>Tag the “firsthello” image as shown below, command is “<code>docker tag firsthello msrinivascharan/docker-kubernetes:module1</code>”</p> <pre>[root@ip-172-42-32-4 DockerWS]# docker tag firsthello msrinivascharan/docker-kubernetes:module1 [root@ip-172-42-32-4 DockerWS]# docker images REPOSITORY          TAG           IMAGE ID            CREATED             SIZE msrinivascharan/docker-kubernetes    module1        0b25d2c5ab0d        13 minutes ago     131MB firsthello          latest        0b25d2c5ab0d        13 minutes ago     131MB python              2.7-slim      4620f1f365b9        6 days ago        120MB hello-world         latest        fce289e99eb9        4 weeks ago       1.84kB [root@ip-172-42-32-4 DockerWS]#</pre>

## Container orchestration with Kubernetes Hands-on Workbook

	<p>You can Upload your tagged image to the repository as shown below, command is “<code>docker push msrinivascharan/docker-kubernetes:module1</code>”</p> <pre>[root@ip-172-42-32-4 DockerWS]# docker push msrinivascharan/docker-kubernetes:module1 The push refers to repository [docker.io/msrinivascharan/docker-kubernetes] 36844df495e1: Pushed 4bb4bb2fdea0: Pushed 31b9a82d7f36: Pushed 4717ad73473a: Mounted from library/python 126818daa940: Mounted from library/python c06903a8f01f: Mounted from library/python 3c816b4ead84: Mounted from library/python module1: digest: sha256:ba5fe78b8cb806f3138b00d17a21ca00baa0611d429b426911feff2d2b633679 size: 1787 [root@ip-172-42-32-4 DockerWS]# ]</pre> <p>Go to <a href="https://hub.docker.com/">https://hub.docker.com/</a>, you can see recently pushed image as shown below</p>  <p><b>Repositories</b></p> <p>Type to filter repositories by name</p> <table border="1"> <thead> <tr> <th>msrinivascharan/docker-kubernetes</th> <th>public</th> <th>0 STARS</th> <th>1 PULLS</th> <th>DETAILS</th> </tr> </thead> </table> <p><b>PUBLIC REPOSITORY</b></p> <p><b>msrinivascharan/docker-kubernetes</b> ☆</p> <p>Last pushed: 6 minutes ago</p> <p>Repo Info Tags Collaborators Webhooks Settings</p> <p>Short Description</p> <p>Short description is empty for this repo.</p> <p>Docker Pull Command</p> <p>docker pull msrinivascharan/docker-kubernetes:module1</p> <p>Full Description</p> <p>Full description is empty for this repo.</p> <p>Owner</p> <p>msrinivascharan</p>	msrinivascharan/docker-kubernetes	public	0 STARS	1 PULLS	DETAILS
msrinivascharan/docker-kubernetes	public	0 STARS	1 PULLS	DETAILS		
11	<p>Once image uploaded to docker hub or any repository, you can pull from any ware and run as shown below. Command is “<code>docker run -p 4000:80 msrinivascharan/docker-kubernetes:module1</code> - on diff machine”</p> <p>Note: If the image isn't available locally on the machine, Docker pulls it from the registry.</p>					

----- End of Exercise -----

***Module 2 – Setting up single node cluster with Kubernetes Minikube*****Exercise 2.1 Setting up single node Kubernetes cluster using Minikube****Scenario**

Kubectl and Minikube installation

**To do tasks:**

In this exercise, you will be:

- Kubectl installation
- Minikube installation

**Steps**

Kubectl & Docker, Minikube setup	
1	Go to directory “/etc/yum.repos.d” and create new file “kubernetes.repo”
2	<p>Add below mentioned content to newly created “kubernetes.repo” file, as shown below</p> <pre>[user1@ip-172-42-0-39 yum.repos.d]\$ sudo cat kubernetes.repo [kubernetes] name=Kubernetes baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-x86_64 enabled=1 gpgcheck=1 repo_gpgcheck=1 gpgkey=https://packages.cloud.google.com/yum/doc/yum-key.gpg https://packages.cloud.google.com/yum/doc/rpm-package-key.gpg [user1@ip-172-42-0-39 yum.repos.d]\$ █</pre> <p>You can copy above mentioned configurations from <a href="#">here</a> or as shown below</p> <pre>[kubernetes] name=Kubernetes baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-x86_64 enabled=1 gpgcheck=1 repo_gpgcheck=1 gpgkey=https://packages.cloud.google.com/yum/doc/yum-key.gpg https://packages.cloud.google.com/yum/doc/rpm-package-key.gpg</pre>
3	Run command as shown below to <b>Install Kubectl</b> command is “yum install -y kubectl”

## Container orchestration with Kubernetes Hands-on Workbook

	<pre>[user1@ip-172-42-0-39 ~]\$ sudo yum install -y kubectl Loaded plugins: fastestmirror, langpacks base epel/x86_64/metalink epel http://del-repos.extreme-ix.org/epel/7/x86_64/repo/repodata/repomd.xml: [Errno -1] repomd.xml does not match metalink for epel Trying other mirror. epel extras kubernetes/signature Retrieving key from https://packages.cloud.google.com/yum/doc/yum-key.gpg Importing GPG key 0xA7317B0F:   Userid : "Google Cloud Packages Automatic Signing Key &lt;gc-team@google.com&gt;"   Fingerprint: d0bc 747f d8ca f711 7500 d6fa 3746 c208 a731 7b0f   From   : https://packages.cloud.google.com/yum/doc/yum-key.gpg Retrieving key from https://packages.cloud.google.com/yum/doc/rpm-package-key.gpg kubernetes/signature nux-dextop updates xrdp (1/10): base/7/x86_64/group_gz (2/10): extras/7/x86_64/primary_db (3/10): kubernetes/primary (4/10): epel/x86_64/group_gz (5/10): updates/7/x86_64/primary_db (6/10): nux-dextop/x86_64/primary_db (7/10): base/7/x86_64/primary_db (8/10): epel/x86_64/primary (9/10): epel/x86_64/updateinfo (10/10): xrdp/primary_db epel kubernetes Resolving Dependencies --&gt; Running transaction check --&gt; Package kubectl.x86_64 0:1.11.0-0 will be installed --&gt; Finished Dependency Resolution  Dependencies Resolved  ===== Package           Arch      Version       Repository      Size ===== Installing:   kubectl          x86_64    1.11.0-0     kubernetes    7.5 M  Transaction Summary ===== Install 1 Package  Total download size: 7.5 M Installed size: 37 M Downloading packages: warning: /var/cache/yum/x86_64/7/kubernetes/packages/5736d31b18c9a00419105394d462ecec4847e6a9bbc7b5a4e41790a67e29c817-kubectl-1.11.0-0.x86_64.rpm: Header V4 RSA/SHA512 Signature, key ID 3e1ba8d5: NOKEY Public key for 5736d31b18c9a00419105394d462ecec4847e6a9bbc7b5a4e41790a67e29c817-kubectl-1.11.0-0.x86_64.rpm is not installed 5736d31b18c9a00419105394d462ecec4847e6a9bbc7b5a4e41790a67e29c817-kubectl-1.11.0-0.x86_64.rpm   7.5 MB  00:00:01 Retrieving key from https://packages.cloud.google.com/yum/doc/yum-key.gpg Importing GPG key 0xA7317B0F:   Userid : "Google Cloud Packages Automatic Signing Key &lt;gc-team@google.com&gt;"   Fingerprint: d0bc 747f d8ca f711 7500 d6fa 3746 c208 a731 7b0f   From   : https://packages.cloud.google.com/yum/doc/yum-key.gpg Retrieving key from https://packages.cloud.google.com/yum/doc/rpm-package-key.gpg Importing GPG key 0x3E1BA8D5:   Userid : "Google Cloud Packages RPM Signing Key &lt;gc-team@google.com&gt;"   Fingerprint: 3749 elba 95a8 6ce0 5454 6ed2 f09c 394c 3e1b a8d5   From   : https://packages.cloud.google.com/yum/doc/rpm-package-key.gpg Running transaction check Running transaction test Transaction test succeeded Running transaction   Installing : kubectl-1.11.0-0.x86_64   Verifying  : kubectl-1.11.0-0.x86_64   Installed: kubectl.x86_64 0:1.11.0-0   Complete!   1/1 1/1  Run command "kubectl" to check the installation </pre>
4	<p>Note: Docker installation mandatory before proceeding with Minikube installation</p> <p>Use below steps to install Minikube on Linux (CentOS) environment</p> <p>Step 1: cd to downloads -&gt; curl -Lo minikube <a href="https://storage.googleapis.com/minikube/releases/latest/minikube-linux-amd64">https://storage.googleapis.com/minikube/releases/latest/minikube-linux-amd64</a></p>

## Container orchestration with Kubernetes Hands-on Workbook

	<pre>[root@ip-172-42-32-4 Downloads]# ls minikube [root@ip-172-42-32-4 Downloads]#</pre> <p>Step 2: sudo mv minikube /usr/bin/</p> <p>Step 3: sudo chmod +x minikube</p> <pre>[root@ip-172-42-32-4 Downloads]# sudo chmod +x minikube [root@ip-172-42-32-4 Downloads]# ls minikube [root@ip-172-42-32-4 Downloads]# sudo mv minikube /usr/bin/ [root@ip-172-42-32-4 Downloads]# ls [root@ip-172-42-32-4 Downloads]#</pre>
5	<p>Use below steps to start Minikube on Linux (CentOS) environment, command is “<code>minikube start --vm-driver=none</code>”</p> <pre>[root@ip-172-42-32-4 /]# minikube start --vm-driver=none Starting local Kubernetes v1.13.2 cluster... Starting VM... Getting VM IP address... Moving files into cluster... Downloading kubelet v1.13.2 Downloading kubelet v1.13.2 Finished Downloading kubelet v1.13.2 Finished Downloading kubelet v1.13.2 Setting up certs... Connecting to cluster... Setting up kubeconfig... Stopping extra container runtimes... Starting cluster components... Verifying kubelet health ... Verifying apiserver health ... Kubectl is now configured to use the cluster. ===== WARNING: IT IS RECOMMENDED NOT TO RUN THE NONE DRIVER ON PERSONAL WORKSTATIONS The 'none' driver will run an insecure kubernetes apiserver as root that may leave the host vulnerable to CSRF attacks</pre> <p>When using the none driver, the kubectl config and credentials generated will be root owned and will appear in the root home directory. You will need to move the files to the appropriate location and then set the correct permissions. An example of this is below:</p> <pre>sudo mv /root/.kube \$HOME/.kube # this will write over any previous configuration sudo chown -R \$USER \$HOME/.kube sudo chgrp -R \$USER \$HOME/.kube  sudo mv /root/.minikube \$HOME/.minikube # this will write over any previous configuration sudo chown -R \$USER \$HOME/.minikube sudo chgrp -R \$USER \$HOME/.minikube</pre> <p>This can also be done automatically by setting the env var <code>CHANGE_MINIKUBE_NONE_USER=true</code></p> <p>Loading cached images from config file.</p> <p>Everything looks great. Please enjoy minikube!</p> <pre>[root@ip-172-42-32-4 /]#</pre>
6	<p>Run command “” as shown below to Check that kubectl is properly configured by getting the cluster state</p> <pre>[user1@ip-172-42-0-39 /]\$ kubectl cluster-info Kubernetes master is running at https://172.42.0.39:8443 KubeDNS is running at https://172.42.0.39:8443/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy  To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'. [user1@ip-172-42-0-39 /]\$</pre>
7	

## Container orchestration with Kubernetes Hands-on Workbook

Run command “kubectl get nodes” as shown below to see available node details

```
[root@ip-172-42-32-4 /]# kubectl get nodes
NAME      STATUS  ROLES   AGE    VERSION
minikube  Ready   master  113m   v1.13.2
[root@ip-172-42-32-4 /]#
```

----- End of Exercise -----

---

### Module 3 – Kubernetes Cluster basic

---

## Exercise 3.1 Deploying an App to Kubernetes cluster

### Scenario

Deploy first app on to Kubernetes cluster

### To do tasks:

In this exercise, you will be:

- Deploy first app on Kubernetes using Kubectl
- Basics of Kubectl and its CLI, and App interactions

### Steps

#### Deploying app

1

Check nodes availability as shown below using command “`sudo kubectl get nodes`”

```
[root@ip-172-42-32-4 /]# kubectl get nodes
NAME      STATUS  ROLES   AGE    VERSION
minikube  Ready   master  113m   v1.13.2
[root@ip-172-42-32-4 /]#
```

2

Run command “`kubectl run kubernetes-bootcamp --image=gcr.io/google-samples/kubernetes-bootcamp:v1 --port=8080`” as shown below to deploy the app

```
[user1@ip-172-42-32-7 /]$ sudo kubectl run kubernetes-bootcamp --image=gcr.io/google-samples/kubernetes-bootcamp:v1 --port=8080
deployment.apps/kubernetes-bootcamp created
[user1@ip-172-42-32-7 /]$
```

The run command creates a new deployment. We need to provide the deployment name and app image location (include the full repository URL for images hosted outside Docker hub). We want to run the app on a specific port, so we add the `--port` parameter

## Container orchestration with Kubernetes Hands-on Workbook

	<p>Above command done 3 things</p> <ul style="list-style-type: none"> <li>• searched for a suitable node where an instance of the application could be run (we have only 1 available node)</li> <li>• scheduled the application to run on that Node</li> <li>• configured the cluster to reschedule the instance on a new Node when needed</li> </ul>
3	<p>Run command “<code>kubectl get deployments</code>” as shown below to get deployed app details</p> <pre>[root@ip-172-42-32-4 /]# kubectl get deployments NAME          READY   UP-TO-DATE   AVAILABLE   AGE kubernetes-bootcamp   1/1      1           1          98s [root@ip-172-42-32-4 /]#</pre> <p>We see that there is 1 deployment running a single instance of your app. The instance is running inside a Docker container on your node</p>
4	<p>Run command “<code>kubectl proxy</code>” as shown below in terminal 2 to run proxy</p> <pre>[user1@ip-172-42-32-7 ~]\$ sudo kubectl proxy [sudo] password for user1: Starting to serve on 127.0.0.1:8001</pre> <p>Note: The <code>kubectl</code> command can create a proxy that will forward communications into the cluster-wide, private network. The proxy can be terminated by pressing control-C and won't show any output while its running.</p>
5	<p>Run command “<code>curl http://localhost:8001/version</code>” as shown below in terminal to check proxy</p> <pre>[root@ip-172-42-32-4 /]# curl http://localhost:8001/version {   "major": "1",   "minor": "13",   "gitVersion": "v1.13.2",   "gitCommit": "cff46ab41ff0bb44d8584413b598ad8360ec1def",   "gitTreeState": "clean",   "buildDate": "2019-01-10T23:28:14Z",   "goVersion": "go1.11.4",   "compiler": "gc",   "platform": "linux/amd64" }[root@ip-172-42-32-4 /]#</pre>
6	<p>A Pod represents a set of running containers on your cluster</p> <p>Run command “<code>kubectl get pods</code>”</p>

## Container orchestration with Kubernetes Hands-on Workbook

	<pre>[root@ip-172-42-32-4 /]# kubectl get pods NAME                               READY   STATUS    RESTARTS   AGE kubernetes-bootcamp-6bf84cb898-dgk8f   1/1     Running   0          11m [root@ip-172-42-32-4 /]#</pre>
7	<p>Run command “curl http://localhost:8001/api/v1/namespaces/default/pods/&lt; kubernetes-bootcamp-6bf84cb898-dgk8f &gt;/proxy/” to access app UI</p> <pre>[root@ip-172-42-32-4 /]# curl http://localhost:8001/api/v1/namespaces/default/pods/kubernetes-bootcamp-6bf84cb898-dgk8f/proxy/ Hello Kubernetes bootcamp!   Running on: kubernetes-bootcamp-6bf84cb898-dgk8f   v=1 [root@ip-172-42-32-4 /]#</pre>

----- End of Exercise -----

## Exercise 3.2 Exploring App Deployed in Kubernetes cluster

## Scenario

Exploring deployed app using kubectl

## To do tasks:

In this exercise, you will be:

- Troubleshoot Kubernetes application using kubectl
- Using kubectl get, describe, log, exec commands

## Steps

Exploring App	
1	<p>Run command “sudo kubectl get pod” as shown below to get deployed app details</p> <pre>[root@ip-172-42-32-4 /]# kubectl get pods NAME                               READY   STATUS    RESTARTS   AGE kubernetes-bootcamp-6bf84cb898-dgk8f   1/1     Running   0          11m [root@ip-172-42-32-4 /]#</pre>
2	<p>Run command “kubectl describe pods” as shown below to get more information about pods</p>

## Container orchestration with Kubernetes Hands-on Workbook

	<pre>[root@ip-172-42-32-4 /]# kubectl describe pods Name:           kubernetes-bootcamp-6bf84cb898-dgk8f Namespace:      default Priority:       0 PriorityClassName: &lt;none&gt; Node:           minikube/172.42.32.4 Start Time:     Fri, 01 Feb 2019 11:18:05 +0000 Labels:         pod-template-hash=6bf84cb898                 run=kubernetes-bootcamp Annotations:   &lt;none&gt; Status:        Running IP:            172.17.0.4 Controlled By: ReplicaSet/kubernetes-bootcamp-6bf84cb898 Containers:   kubernetes-bootcamp:     Container ID: docker://f04126d00b98164299dcefc0c54a5adc6d025d5f0ff96fef3c1c2f9cb51261     Image:          gcr.io/google-samples/kubernetes-bootcamp:v1     Image ID:      docker-pullable://gcr.io/google-samples/kubernetes-bootcamp@sha256:0d6b8ee63bb57c5f5b6156f446b3bc3b3c143d233037f3a2f00e279c8fcc64af     Port:          8080/TCP     Host Port:    0/TCP     State:        Running     Started:     Fri, 01 Feb 2019 11:18:16 +0000     Ready:        True     Restart Count: 0     Environment:  &lt;none&gt;     Mounts:       /var/run/secrets/kubernetes.io/serviceaccount from default-token-rmfdk (ro) Conditions:   Type  Status   Initialized  True   Ready  True   ContainersReady  True   PodScheduled  True Volumes:   default-token-rmfdk:     Type:  Secret (a volume populated by a Secret)     SecretName: default-token-rmfdk     Optional:  false QoS Class:  BestEffort Node-Selectors: &lt;none&gt; Tolerations:  node.kubernetes.io/not-ready:NoExecute for 300s               node.kubernetes.io/unreachable:NoExecute for 300s Events:   Type  Reason  Age From      Message   ----  ----  --  ----      -----   Normal Scheduled  23m default-scheduler  Successfully assigned default/kubernetes-bootcamp-6bf84cb898-dgk8f to minikube   Normal Pulling   23m kubelet, minikube  pulling image "gcr.io/google-samples/kubernetes-bootcamp:v1"   Normal Pulled    22m kubelet, minikube  Successfully pulled image "gcr.io/google-samples/kubernetes-bootcamp:v1"   Normal Created   22m kubelet, minikube  Created container   Normal Started   22m kubelet, minikube  Started container [root@ip-172-42-32-4 /]#</pre> <p>Note: the describe command can be used to get detailed information about most of the kubernetes primitives: node, pods, deployments. The describe output is designed to be human readable, not to be scripted against.</p>
3	<p>Run command as shown below in terminal 2 to run proxy</p> <pre>[user1@ip-172-42-32-7 ~]\$ sudo kubectl proxy [sudo] password for user1: Starting to serve on 127.0.0.1:8001</pre> <p>Note: The kubectl command can create a proxy that will forward communications into the cluster-wide, private network. The proxy can be terminated by pressing control-C and won't show any output while its running.</p>
4	<p>Run command “kubectl logs &lt; kubernetes-bootcamp-6bf84cb898-dgk8f &gt;” as shown below to see logs</p> <pre>[root@ip-172-42-32-4 /]# kubectl logs kubernetes-bootcamp-6bf84cb898-dgk8f Kubernetes Bootcamp App Started At: 2019-02-01T11:18:16.704Z   Running On: kubernetes-bootcamp-6bf84cb898-dgk8f Running On: kubernetes-bootcamp-6bf84cb898-dgk8f   Total Requests: 1   App Uptime: 1016.742 seconds   Log Time: 2019-02-01T11:35:13.446Z Running On: kubernetes-bootcamp-6bf84cb898-dgk8f   Total Requests: 2   App Uptime: 1257.826 seconds   Log Time: 2019-02-01T11:39:14.530Z [root@ip-172-42-32-4 /]#</pre>
6	

## Container orchestration with Kubernetes Hands-on Workbook

	<p>Run command “kubectl exec &lt; kubernetes-bootcamp-6bf84cb898-dgk8f &gt; env” as shown below to execute commands directly on the container to list the environment variables</p> <pre>[root@ip-172-42-32-4 /]# kubectl exec kubernetes-bootcamp-6bf84cb898-dgk8f env PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin HOSTNAME=kubernetes-bootcamp-6bf84cb898-dgk8f KUBERNETES_PORT_443_TCP=tcp://10.96.0.1:443 KUBERNETES_PORT_443_TCP_PROTO=tcp KUBERNETES_PORT_443_TCP_PORT=443 KUBERNETES_PORT_443_TCP_ADDR=10.96.0.1 KUBERNETES_SERVICE_HOST=10.96.0.1 KUBERNETES_SERVICE_PORT=443 KUBERNETES_SERVICE_PORT_HTTPS=443 KUBERNETES_PORT=tcp://10.96.0.1:443 NPM_CONFIG_LOGLEVEL=info NODE_VERSION=6.3.1 HOME=/root [root@ip-172-42-32-4 /]# ]</pre> <p>Note: name of the container omitted in the above command, since we only have a single container in the Pod.</p>
7	<p>Run command “sudo kubectl exec -ti &lt; kubernetes-bootcamp-6bf84cb898-dgk8f &gt; bash” as shown below to open console on the container where we run our application (NodeJS).</p> <pre>[root@ip-172-42-32-4 /]# sudo kubectl exec -ti kubernetes-bootcamp-6bf84cb898-dgk8f bash root@kubernetes-bootcamp-6bf84cb898-dgk8f:/# ]</pre>
8	<p>Run command as shown below to see source code of the containerized app, code is in the server.js file</p> <pre>[root@ip-172-42-32-4 /]# sudo kubectl exec -ti kubernetes-bootcamp-6bf84cb898-dgk8f bash root@kubernetes-bootcamp-6bf84cb898-dgk8f:/# cat server.js var http = require('http'); var requests=0; var podname= process.env.HOSTNAME; var startTime; var host; var handleRequest = function(request, response) {   response.setHeader('Content-Type', 'text/plain');   response.writeHead(200);   response.write("Hello Kubernetes bootcamp!   Running on: ");   response.write(host);   response.end("   v=1\n");   console.log("Running On:" ,host, "  Total Requests:", ++requests,"  App Uptime:", (new Date() - startTime)/1000 , "seconds", "  Log Time:",new Date()); } var www = http.createServer(handleRequest); www.listen(8080,function () {   startTime = new Date();   host = process.env.HOSTNAME;   console.log ("Kubernetes Bootcamp App Started At:",startTime, "  Running On: " ,host, "\n"); }); root@kubernetes-bootcamp-6bf84cb898-dgk8f:/# ]</pre>
	<p>Run command “curl localhost:8080” as shown below to check that the application is up</p> <pre>root@kubernetes-bootcamp-6bf84cb898-dgk8f:/# curl localhost:8080 Hello Kubernetes bootcamp!   Running on: kubernetes-bootcamp-6bf84cb898-dgk8f   v=1 root@kubernetes-bootcamp-6bf84cb898-dgk8f:/# ]</pre> <p>Note: here we used localhost because we executed the command inside the NodeJS container. In exercise 1.2 step 7, we did by reaching pod through proxy</p>

## Container orchestration with Kubernetes Hands-on Workbook

You can run “node -v (NodeJS command)” and “node -p ‘3+2’” commands also

----- End of Exercise -----

### Exercise 3.3 Exposing deployed app using services

#### Scenario

Exposing deployed app using service

#### To do tasks:

In this exercise, you will be:

- Expose Kubernetes app outside the cluster using kubectl expose command
- View and apply labels to object using kubectl label command

#### Steps

##### Exposing App

1

Run command “kubectl get pods,services” as shown below to get existing pods and services

```
[root@ip-172-42-32-4 /]# kubectl get pods,services
NAME                               READY   STATUS    RESTARTS   AGE
pod/kubernetes-bootcamp-6bf84cb898-dgk8f   1/1     Running   0          62m

NAME            TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
service/kubernetes   ClusterIP  10.96.0.1   <none>        443/TCP   178m
[root@ip-172-42-32-4 /]#
```

2

Run command “kubectl expose deployment/kubernetes-bootcamp --type=”NodePort” --port 8080” as shown below to create a service

```
[root@ip-172-42-32-4 /]# kubectl expose deployment/kubernetes-bootcamp --type="NodePort" --port 8080
service/kubernetes-bootcamp exposed
[root@ip-172-42-32-4 /]#
```

You can check newly created sevices using get command “kubectl get services” as shown below

```
[root@ip-172-42-32-4 /]# kubectl get services
NAME            TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
kubernetes     ClusterIP  10.96.0.1   <none>        443/TCP   3h1m
kubernetes-bootcamp   NodePort   10.100.178.89   <none>        8080:30119/TCP   67s
[root@ip-172-42-32-4 /]#
```

We have now a running Service called **kubernetes-bootcamp**. Here we see that the Service received a unique cluster-IP, an internal port and an external-IP (the IP of the Node).

## Container orchestration with Kubernetes Hands-on Workbook

3	<p>Run command “<code>kubectl describe service/kubernetes-bootcamp</code>” as shown below to see what port was opened externally</p> <pre>[root@ip-172-42-32-4 /]# kubectl describe service/kubernetes-bootcamp Name:           kubernetes-bootcamp Namespace:      default Labels:         run=kubernetes-bootcamp Annotations:   &lt;none&gt; Selector:       run=kubernetes-bootcamp Type:          NodePort IP:            10.100.178.89 Port:          &lt;unset&gt;  8080/TCP TargetPort:    8080/TCP NodePort:      &lt;unset&gt;  30119/TCP Endpoints:    172.17.0.4:8080 Session Affinity: None External Traffic Policy: Cluster Events:        &lt;none&gt; [root@ip-172-42-32-4 /]#</pre>
4	<p>Try below mentioned options to access app UI</p> <ol style="list-style-type: none"> <li>1]curl <a href="http://&lt;service IP&gt;">http://&lt;service IP&gt;</a>: target port</li> <li>2]curl <a href="http://&lt;pod IP&gt;">http://&lt;pod IP&gt;</a>:target port</li> <li>3]curl <a href="http://&lt;node IP&gt;">http://&lt;node IP&gt;</a>:NodePort</li> </ol> <p>Note: try all combinations of IP and ports, out side cluster also.</p>
6	<p>Run command “<code>kubectl describe service/kubernetes-bootcamp</code>” as shown below to see label created by default</p> <pre>[root@ip-172-42-32-4 /]# kubectl describe service/kubernetes-bootcamp Name:           kubernetes-bootcamp Namespace:      default Labels:         run=kubernetes-bootcamp Annotations:   &lt;none&gt; Selector:       run=kubernetes-bootcamp Type:          NodePort IP:            10.100.178.89 Port:          &lt;unset&gt;  8080/TCP TargetPort:    8080/TCP NodePort:      &lt;unset&gt;  30119/TCP Endpoints:    172.17.0.4:8080 Session Affinity: None External Traffic Policy: Cluster Events:        &lt;none&gt; [root@ip-172-42-32-4 /]#</pre>
7	<p>Run command “<code>kubectl get pods -l run=kubernetes-bootcamp</code>” as shown below to get pods associated with label</p>

## Container orchestration with Kubernetes Hands-on Workbook

	<pre>[root@ip-172-42-32-4 /]# kubectl get pods -l run=kubernetes-bootcamp NAME                  READY   STATUS    RESTARTS   AGE kubernetes-bootcamp-6bf84cb898-dgk8f   1/1     Running   0          80m [root@ip-172-42-32-4 /]#</pre>
9	<p>Run command as shown below to <a href="#">add label to pod</a>.</p> <p>To apply a new label, we use the label command followed by the object type, object name and the new label</p> <pre>"kubectl label pod &lt;kubernetes-bootcamp-6bf84cb898-dgk8f&gt; app=v1"</pre> <pre>[root@ip-172-42-32-4 /]# kubectl label pod kubernetes-bootcamp-6bf84cb898-dgk8f app=v1 pod/kubernetes-bootcamp-6bf84cb898-dgk8f labeled [root@ip-172-42-32-4 /]#</pre> <p>This will apply a new label app=v1 to our Pod (we pinned the application version to the Pod), and we can check it with the describe pod command.</p>
10	<p>Run command <a href="#">"kubectl get pods -l app=v1"</a> as shown below to see list of pods using the new label</p> <pre>[root@ip-172-42-32-4 /]# kubectl get pods -l app=v1 NAME                  READY   STATUS    RESTARTS   AGE kubernetes-bootcamp-6bf84cb898-dgk8f   1/1     Running   0          85m [root@ip-172-42-32-4 /]#</pre>
11	<p>Run command <a href="#">"kubectl delete service -l run=kubernetes-bootcamp"</a> as shown below to delete service</p> <pre>[user1@ip-172-42-32-7 ~]\$ sudo kubectl get services NAME            TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)        AGE kubernetes      ClusterIP  10.96.0.1      &lt;none&gt;           443/TCP       4h kubernetes-bootcamp  NodePort   10.101.211.56  &lt;none&gt;           8080:32487/TCP  1h [user1@ip-172-42-32-7 ~]\$ sudo kubectl delete service -l run=kubernetes-bootcamp service "kubernetes-bootcamp" deleted [user1@ip-172-42-32-7 ~]\$ sudo kubectl get services NAME            TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)        AGE kubernetes      ClusterIP  10.96.0.1      &lt;none&gt;           443/TCP       4h [user1@ip-172-42-32-7 ~]\$</pre> <p>This confirms that our Service was removed. To confirm that route is not exposed anymore you can curl the previously exposed IP and port</p> <p>Above command proves that the app is not reachable anymore from outside of the cluster</p> <p style="text-align: center;">----- End of Exercise -----</p>

## Exercise 3.4 Scaling your App

### Scenario

Scale a deployment with Kubectl

## Container orchestration with Kubernetes Hands-on Workbook

## To do tasks:

In this exercise, you will be:

- Scale deployment with kubectl scale
- Load balancing

## Steps

## Scaling app

1	<p>Run command “<code>kubectl get deployments</code>” as shown below to check deployments</p> <pre>[root@ip-172-42-32-4 /]# kubectl get deployments NAME          READY   UP-TO-DATE   AVAILABLE   AGE kubernetes-bootcamp  1/1      1           1          93m [root@ip-172-42-32-4 /]#</pre> <p>We should have 1 Pod. If not, run the command again.</p> <p>The <b>DESIRED</b> state is showing the configured number of replicas</p> <p>The <b>CURRENT</b> state shows how many replicas are running now</p> <p>The <b>UP-TO-DATE</b> is the number of replicas that were updated to match the desired (configured) state</p> <p>The <b>AVAILABLE</b> state shows how many replicas are AVAILABLE to the users</p>
2	<p>Run command as shown below to scale the Deployment to 4 replicas. Use the <code>kubectl scale</code> command, followed by the deployment type, name and desired number of instances</p> <p>“<code>kubectl scale deployments/kubernetes-bootcamp --replicas=4</code>”</p> <pre>[root@ip-172-42-32-4 /]# kubectl scale deployments/kubernetes-bootcamp --replicas=4 deployment.extensions/kubernetes-bootcamp scaled [root@ip-172-42-32-4 /]#</pre> <p>To list your Deployments once again, run command “<code>kubectl get deployments</code>” as shown below</p> <pre>[root@ip-172-42-32-4 /]# kubectl get deployments NAME          READY   UP-TO-DATE   AVAILABLE   AGE kubernetes-bootcamp  4/4      4           4          98m [root@ip-172-42-32-4 /]#</pre>
3	<p>We have 4 instances of the application available. Check if the number of Pods changed using command “<code>kubectl get pods -o wide</code>”</p> <pre>[root@ip-172-42-32-4 /]# kubectl get pods -o wide NAME          READY   STATUS    RESTARTS   AGE     IP            NODE   NOMINATED NODE   READINESS GATES kubernetes-bootcamp-6bf84cb898-524dq  1/1   Running   0          2m26s  172.17.0.5  minikube  &lt;none&gt;        &lt;none&gt; kubernetes-bootcamp-6bf84cb898-dgk8f  1/1   Running   0          99m    172.17.0.4  minikube  &lt;none&gt;        &lt;none&gt; kubernetes-bootcamp-6bf84cb898-s4ps4  1/1   Running   0          2m26s  172.17.0.7  minikube  &lt;none&gt;        &lt;none&gt; kubernetes-bootcamp-6bf84cb898-wd6m7  1/1   Running   0          2m26s  172.17.0.6  minikube  &lt;none&gt;        &lt;none&gt; [root@ip-172-42-32-4 /]#</pre>

## Container orchestration with Kubernetes Hands-on Workbook

	<p>Note: -o=wide id for Output in the plain-text format with any additional information, and for pods, the node name is included</p>
4	<p>The change was registered in the Deployment events log. To check that, use the describe command “<code>kubectl describe deployments/kubernetes-bootcamp</code>”</p> <pre>[root@ip-172-42-32-4 /]# kubectl describe deployments/kubernetes-bootcamp Name:           kubernetes-bootcamp Namespace:      default CreationTimestamp: Fri, 01 Feb 2019 11:18:05 +0000 Labels:         run=kubernetes-bootcamp Annotations:   deployment.kubernetes.io/revision: 1 Selector:       run=kubernetes-bootcamp Replicas:      4 desired   4 updated   4 total   4 available   0 unavailable StrategyType:  RollingUpdate MinReadySeconds: 0 RollingUpdateStrategy: 25% max unavailable, 25% max surge Pod Template:   Labels:  run=kubernetes-bootcamp   Containers:     kubernetes-bootcamp:       Image:      gcr.io/google-samples/kubernetes-bootcamp:v1       Port:       8080/TCP       Host Port:  0/TCP       Environment: &lt;none&gt;       Mounts:    &lt;none&gt;       Volumes:   &lt;none&gt;   Conditions:     Type     Status  Reason     ----  -----     Progressing  True   NewReplicaSetAvailable     Available   True   MinimumReplicasAvailable     OldReplicaSets: &lt;none&gt;     NewReplicaSet:  kubernetes-bootcamp-6bf84cb898 (4/4 replicas created)   Events:     Type     Reason          Age   From            Message     ----  -----     Normal  ScalingReplicaSet 3m11s  deployment-controller  Scaled up replica set kubernetes-bootcamp-6bf84cb898 to 4 [root@ip-172-42-32-4 /]# NewReplicaSet:  kubernetes-bootcamp-5c69669756 (4/4 replicas created) Events:   Type     Reason          Age   From            Message   ----  -----   Normal  ScalingReplicaSet 17m   deployment-controller  Scaled up replica set kubernetes-bootcamp-5c69669756 to 4 [user1@ip-172-42-32-7 ~]\$ </pre>
5	<p>Run command “<code>kubectl expose deployment/kubernetes-bootcamp --type="NodePort" --port 8080</code>” as shown below to create a service</p> <pre>[user1@ip-172-42-32-7 ~]\$ sudo kubectl expose deployment/kubernetes-bootcamp --type="NodePort" --port 8080 service/kubernetes-bootcamp exposed [user1@ip-172-42-32-7 ~]\$ </pre> <p><b>Note:</b> service creation not required if you didn't disturb (deleted) service created in exercise 3.3 step for the same deployment</p> <p>You can check newly created services using get command “<code>kubectl get services</code>” as shown below</p>

## Container orchestration with Kubernetes Hands-on Workbook

```
[root@ip-172-42-32-4 /]# kubectl get services
NAME           TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE
kubernetes     ClusterIP  10.96.0.1   <none>        443/TCP         3h39m
kubernetes-bootcamp   NodePort   10.100.178.89 <none>        8080:30119/TCP  39m
[root@ip-172-42-32-4 /]# █
```

We have now a running Service called **kubernetes-bootcamp**. Here we see that the Service received a unique cluster-IP, an internal port and an external-IP (the IP of the Node).

"kubectl describe svc kubernetes-bootcamp"

```
[root@ip-172-42-32-4 /]# kubectl describe svc kubernetes-bootcamp
Name:           kubernetes-bootcamp
Namespace:      default
Labels:         run=kubernetes-bootcamp
Annotations:   <none>
Selector:       run=kubernetes-bootcamp
Type:          NodePort
IP:            10.100.178.89
Port:          <unset>  8080/TCP
TargetPort:    8080/TCP
NodePort:      <unset>  30119/TCP
Endpoints:     172.17.0.4:8080,172.17.0.5:8080,172.17.0.6:8080 + 1 more...
Session Affinity: None
External Traffic Policy: Cluster
Events:        <none>
[root@ip-172-42-32-4 /]# █
```

Try access app UI with service IP and port to see load balancer working, We must send multiple requests to see load balancer working

Container orchestration with Kubernetes Hands-on Workbook

7	<p>Run command “<code>kubectl scale deployments/kubernetes-bootcamp --replicas=2</code>” as shown below to scale down the deployments</p> <hr/> <pre>[root@ip-172-42-32-4 /]# kubectl scale deployments/kubernetes-bootcamp --replicas=2 deployment.extensions/kubernetes-bootcamp scaled [root@ip-172-42-32-4 /]#</pre> <p>You can check scale in as shown below</p>
---	---

## Container orchestration with Kubernetes Hands-on Workbook

```
[root@ip-172-42-32-4 /]# kubectl get pods,deployments
NAME                               READY   STATUS    RESTARTS   AGE
pod/kubernetes-bootcamp-6bf84cb898-dgk8f   1/1     Running   0          111m
pod/kubernetes-bootcamp-6bf84cb898-wd6m7   1/1     Running   0          14m

NAME                           READY   UP-TO-DATE   AVAILABLE   AGE
deployment.extensions/kubernetes-bootcamp   2/2      2           2          111m
[root@ip-172-42-32-4 /]#
```

----- End of Exercise -----

## Exercise 3.5 Updating App

### Scenario

Update a deployed application

### To do tasks:

In this exercise, you will be:

- Updating and rollback app using kubectl set and rollout commands

### Steps

Updating App	
2	<p>Run commands as shown below to get the deployment and details</p> <pre>[root@ip-172-42-32-4 /]# kubectl get pods,deployments NAME                               READY   STATUS    RESTARTS   AGE pod/kubernetes-bootcamp-6bf84cb898-dgk8f   1/1     Running   0          111m pod/kubernetes-bootcamp-6bf84cb898-wd6m7   1/1     Running   0          14m  NAME                           READY   UP-TO-DATE   AVAILABLE   AGE deployment.extensions/kubernetes-bootcamp   2/2      2           2          111m [root@ip-172-42-32-4 /]#</pre>
3	<p>Run command as shown below to view the current image version of the app, run a describe command against the Pods</p>

## Container orchestration with Kubernetes Hands-on Workbook

	<pre>[root@ip-172-42-32-4 /]# kubectl describe deployments kubernetes-bootcamp Name:           kubernetes-bootcamp Namespace:      default CreationTimestamp: Fri, 01 Feb 2019 11:18:05 +0000 Labels:         run=kubernetes-bootcamp Annotations:    deployment.kubernetes.io/revision: 1 Selector:       run=kubernetes-bootcamp Replicas:       2 desired   2 updated   2 total   2 available   0 unavailable StrategyType:   RollingUpdate MinReadySeconds: 0 RollingUpdateStrategy: 25% max unavailable, 25% max surge Pod Template:   Labels:  run=kubernetes-bootcamp   Containers:     kubernetes-bootcamp:       Image:      gcr.io/google-samples/kubernetes-bootcamp:v1       Port:       8080/TCP       Host Port:  0/TCP       Environment: &lt;none&gt;       Mounts:    &lt;none&gt;       Volumes:   &lt;none&gt;   Conditions:     Type     Status  Reason     ----     ----   -----     Progressing  True   NewReplicaSetAvailable     Available   True   MinimumReplicasAvailable   OldReplicaSets: &lt;none&gt;   NewReplicaSet:  kubernetes-bootcamp-6bf84cb898 (2/2 replicas created)   Events:     Type     Reason          Age   From            Message     ----     ----          --   --   -----     Normal   ScalingReplicaSet 31m   deployment-controller  Scaled up replica set kubernetes-bootcamp-6bf84cb898 to 4     Normal   ScalingReplicaSet 18m   deployment-controller  Scaled down replica set kubernetes-bootcamp-6bf84cb898 to 2 [root@ip-172-42-32-4 /]#</pre>
4	<p>Run command “<code>kubectl set image deployments/kubernetes-bootcamp kubernetes-bootcamp=jocatalin/kubernetes-bootcamp:v2</code>” as shown below to update the image of the application to version 2</p> <pre>[root@ip-172-42-32-4 /]# kubectl set image deployments/kubernetes-bootcamp kubernetes-bootcamp=jocatalin/kubernetes-bootcamp:v2 deployment.extensions/kubernetes-bootcamp image updated [root@ip-172-42-32-4 /]#</pre> <p>The command notified the Deployment to use a different image for the app and initiated a rolling update.</p>
5	<p>Run command as shown below to Check the status of the new Pods, and view the old one terminating</p> <pre>[root@ip-172-42-32-4 /]# kubectl get pods NAME                      READY   STATUS    RESTARTS   AGE kubernetes-bootcamp-5bf4d5689b-2s562  1/1    Running   0          2m47s kubernetes-bootcamp-5bf4d5689b-8g5j8  1/1    Running   0          2m41s [root@ip-172-42-32-4 /]#</pre>

## Container orchestration with Kubernetes Hands-on Workbook

	<pre>[root@ip-172-42-32-4 /]# kubectl describe deployments kubernetes-bootcamp Name: kubernetes-bootcamp Namespace: default CreationTimestamp: Fri, 01 Feb 2019 11:18:05 +0000 Labels: run=kubernetes-bootcamp Annotations: deployment.kubernetes.io/revision: 2 Selector: run=kubernetes-bootcamp Replicas: 2 desired   2 updated   2 total   2 available   0 unavailable StrategyType: RollingUpdate MinReadySeconds: 0 RollingUpdateStrategy: 25% max unavailable, 25% max surge Pod Template:   Labels: run=kubernetes-bootcamp   Containers:     kubernetes-bootcamp:       Image: jocatalin/kubernetes-bootcamp:v2       Port: 8080/TCP       Host Port: 0/TCP       Environment: &lt;none&gt;       Mounts: &lt;none&gt;       Volumes: &lt;none&gt;   Conditions:     Type Status Reason     ----     Available True   MinimumReplicasAvailable     Progressing True   NewReplicaSetAvailable OldReplicaSets: &lt;none&gt; NewReplicaSet: kubernetes-bootcamp-5bf4d5689b (2/2 replicas created) Events:   Type Reason Age From Message   ----   Normal ScalingReplicaSet 39m deployment-controller Scaled up replica set kubernetes-bootcamp-6bf84cb898 to 4   Normal ScalingReplicaSet 26m deployment-controller Scaled down replica set kubernetes-bootcamp-6bf84cb898 to 2   Normal ScalingReplicaSet 4m44s deployment-controller Scaled up replica set kubernetes-bootcamp-5bf4d5689b to 1   Normal ScalingReplicaSet 4m38s deployment-controller Scaled down replica set kubernetes-bootcamp-6bf84cb898 to 1   Normal ScalingReplicaSet 4m38s deployment-controller Scaled up replica set kubernetes-bootcamp-5bf4d5689b to 2   Normal ScalingReplicaSet 4m37s deployment-controller Scaled down replica set kubernetes-bootcamp-6bf84cb898 to 0 [root@ip-172-42-32-4 /]#</pre>
6	Access App UI using service
9	Run command as shown below to see the rollout status

## Container orchestration with Kubernetes Hands-on Workbook

	<pre>[user1@ip-172-42-32-7 ~]\$ sudo kubectl rollout status deployments/kubernetes-bootcamp deployment "kubernetes-bootcamp" successfully rolled out [user1@ip-172-42-32-7 ~]\$</pre>
11	<p>Let's perform another update, and deploy image tagged as v10</p> <pre>"kubectl set image deployments/kubernetes-bootcamp kubernetes-bootcamp=jocatalin/kubernetes-bootcamp:v10"</pre> <pre>[user1@ip-172-42-32-7 ~]\$ sudo kubectl set image deployments/kubernetes-bootcamp kubernetes-bootcamp=jocatalin/kubernetes-bootcamp:v10 deployment.extensions/kubernetes-bootcamp image updated [user1@ip-172-42-32-7 ~]\$</pre>
12	<p>Use get deployments to see the status of the deployments</p> <p>And something is wrong... We do not have the desired number of Pods available. List the Pods again as shown below</p> <pre>[root@ip-172-42-32-4 /]# kubectl get deployments,pods NAME                                READY   UP-TO-DATE   AVAILABLE   AGE deployment.extensions/kubernetes-bootcamp   2/2     1           2          144m  NAME                               READY   STATUS        RESTARTS   AGE pod/kubernetes-bootcamp-5bf4d5689b-2s562   1/1    Running       0          12m pod/kubernetes-bootcamp-5bf4d5689b-8g5j8   1/1    Running       0          12m pod/kubernetes-bootcamp-c46798fc9-wgqzc   0/1    ImagePullBackOff   0          100s [root@ip-172-42-32-4 /]#</pre>
13	<p>Run describe command on the Pods to get more insights</p>

Container orchestration with Kubernetes Hands-on Workbook

	<pre> Started:      Fri, 01 Feb 2019 13:30:36 +0000 Ready:       True Restart Count: 0 Environment: &lt;none&gt; Mounts:     /var/run/secrets/kubernetes.io/serviceaccount from default-token-rmfdk (ro) Conditions: Type          Status Initialized   True Ready         True ContainersReady True PodScheduled  True Volumes: default-token-rmfdk:   Type:           Secret (a volume populated by a Secret)   SecretName:    default-token-rmfdk   Optional:      false QoS Class:    BestEffort Node-Selectors: &lt;none&gt; Tolerations:   node.kubernetes.io/not-ready:NoExecute for 300s                 node.kubernetes.io/unreachable:NoExecute for 300s Events:   Type  Reason  Age   From            Message   ----  -----  ----  --  -----   Normal Scheduled  13m  default-scheduler  Successfully assigned default/kubernetes-bootcamp-5bf4d5689b-8g5j8 to minikube   Normal Pulled   13m  kubelet, minikube  Container image "jocatalin/kubernetes-bootcamp:v2" already present on machine   Normal Created   13m  kubelet, minikube  Created container   Normal Started   13m  kubelet, minikube  Started container  Name:            kubernetes-bootcamp-c46798fc9-wgqzc Namespace:       default Priority:        0 PriorityClassName: &lt;none&gt; Node:            minikube/172.42.32.4 Start Time:     Fri, 01 Feb 2019 13:40:55 +0000 Labels:   pod-template-hash=c46798fc9   run=kubernetes-bootcamp Annotations: Status:          Pending IP:              172.17.0.4 Controlled By:   ReplicaSet/kubernetes-bootcamp-c46798fc9 Containers:   kubernetes-bootcamp:     Container ID:          jocatalin/kubernetes-bootcamp:v10     Image:                 jocatalin/kubernetes-bootcamp:v10     Image ID:               8080/TCP     Host Port:             0/TCP     State:                 Waiting       Reason:               ErrImagePull     Ready:                 False     Restart Count:          0     Environment:           &lt;none&gt;     Mounts:       /var/run/secrets/kubernetes.io/serviceaccount from default-token-rmfdk (ro) Conditions: Type          Status Initialized   True </pre> <p>Issue with v10 in the repository. Let's roll back to our previously working version. We'll use the rollout undo command</p>
14	Run command “ <code>kubectl rollout undo deployments/kubernetes-bootcamp</code> ” as show below to revert the deployment to the previous known state (v2 of the image).
	<pre>[root@ip-172-42-32-4 user1]# kubectl rollout undo deployments/kubernetes-bootcamp deployment.extensions/kubernetes-bootcamp rolled back [root@ip-172-42-32-4 user1]#</pre>
15	To see everything ok after rollback and Pods are running. Check again the image deployed on them as shown below

## Container orchestration with Kubernetes Hands-on Workbook

	<pre>[root@ip-172-42-32-4 /]# kubectl get pods NAME                               READY   STATUS    RESTARTS   AGE kubernetes-bootcamp-5bf4d5689b-2s562  1/1    Running   1          23m kubernetes-bootcamp-5bf4d5689b-8g5j8  1/1    Running   1          23m [root@ip-172-42-32-4 /]#</pre>
16	<p>You can see deployment is using a stable version of the app (v2). The Rollback was successful.</p> <pre>[root@ip-172-42-32-4 /]# kubectl get deployments NAME        READY   UP-TO-DATE   AVAILABLE   AGE kubernetes-bootcamp  2/2     2           2           156m [root@ip-172-42-32-4 /]# kubectl describe deployments kubernetes-bootcamp Name:           kubernetes-bootcamp Namespace:      default CreationTimestamp: Fri, 01 Feb 2019 11:18:05 +0000 Labels:         run=kubernetes-bootcamp Annotations:   deployment.kubernetes.io/revision: 4 Selector:       run=kubernetes-bootcamp Replicas:      2 desired   2 updated   2 total   2 available   0 unavailable StrategyType:  RollingUpdate MinReadySeconds: 0 RollingUpdateStrategy: 25% max unavailable, 25% max surge Pod Template:   Labels:  run=kubernetes-bootcamp   Containers:     kubernetes-bootcamp:       Image:      jocatalin/kubernetes-bootcamp:v2       Port:       8080/TCP       Host Port:  0/TCP       Environment: &lt;none&gt;       Mounts:    &lt;none&gt;       Volumes:   &lt;none&gt;   Conditions:     Type        Status  Reason     ----        ----   -----     Available   True    MinimumReplicasAvailable     Progressing True    NewReplicaSetAvailable OldReplicaSets: &lt;none&gt; NewReplicaSet:  kubernetes-bootcamp-5bf4d5689b (2/2 replicas created) Events:   Type      Reason     Age   From            Message   ----      ----     --   --   -----   Normal    ScalingReplicaSet  60m   deployment-controller  Scaled up replica set kubernetes-bootcamp-6bf84cb898 to 4   Normal    ScalingReplicaSet  47m   deployment-controller  Scaled down replica set kubernetes-bootcamp-6bf84cb898 to 2   Normal    ScalingReplicaSet  25m   deployment-controller  Scaled up replica set kubernetes-bootcamp-5bf4d5689b to 1   Normal    ScalingReplicaSet  25m   deployment-controller  Scaled down replica set kubernetes-bootcamp-6bf84cb898 to 1   Normal    ScalingReplicaSet  25m   deployment-controller  Scaled up replica set kubernetes-bootcamp-5bf4d5689b to 2   Normal    ScalingReplicaSet  25m   deployment-controller  Scaled down replica set kubernetes-bootcamp-6bf84cb898 to 0   Normal    ScalingReplicaSet  14m   deployment-controller  Scaled up replica set kubernetes-bootcamp-c46798fc9 to 1   Normal    ScalingReplicaSet  2m36s  deployment-controller  Scaled down replica set kubernetes-bootcamp-c46798fc9 to 0 [root@ip-172-42-32-4 /]#</pre>

----- End of Exercise -----

**Module 4 – Kubernetes Cluster Advanced**

====

## Container orchestration with Kubernetes Hands-on Workbook

### Exercise 4.1 Deploying FirstHello app onto Kubernetes cluster

#### Scenario

Deploy simple containerized app on to Kubernetes cluster

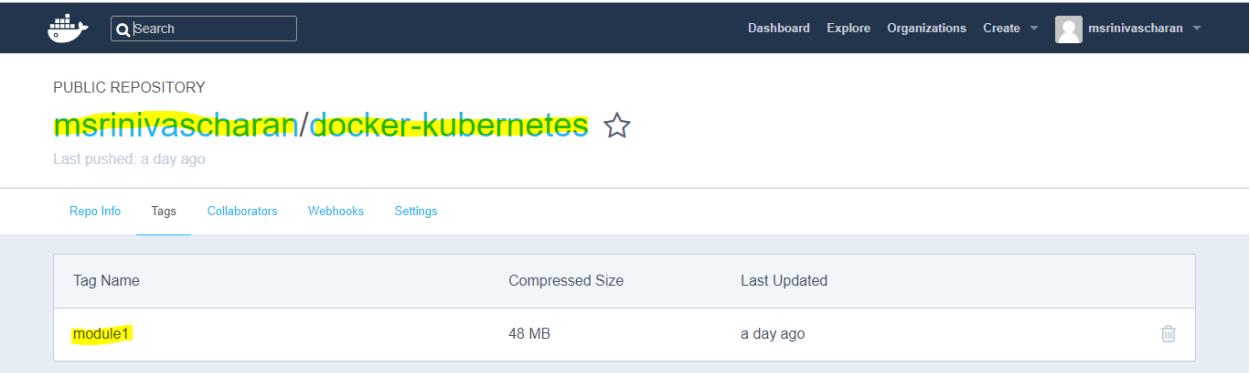
#### To do tasks:

In this exercise, you will be:

- Deploy Python based containerized app on to Kubernetes cluster
- Using Kubernetes replication controller and services

Note: Exercise 1.1 is prerequisite to this exercise

#### Steps

Deploying containerized app on to Kubernetes cluster	
1	Ensure Docker daemon is running
2	Stop and delete “firsthello” app containers if any, use below mentioned commands <ol style="list-style-type: none"><li>1. sudo docker container ls -a</li><li>2. sudo docker stop &lt;containerid&gt;</li><li>3. sudo docker ps -f "status=exited"</li><li>4. sudo rm &lt;containerid&gt;</li></ol>
3	Delete “firsthello” images from local repository if any, use below mentioned commands <ol style="list-style-type: none"><li>1. sudo docker images</li><li>2. sudo docker rmi &lt;imageid&gt;</li></ol>
4	Copy details of “firsthello” image pushed (part of exercise 1.1) to Dockerhub as shown below  <p>The screenshot shows the Dockerhub interface for the repository "msrinivascharan/docker-kubernetes". The repository has 1 star. It was last pushed a day ago. There is one tag listed: "module1". The tag details show a compressed size of 48 MB and was updated a day ago. A trash icon is visible next to the tag row.</p> <p>Create workplace and add deployment and service configurations as shown below</p>

## Container orchestration with Kubernetes Hands-on Workbook

	<pre>[user1@ip-172-42-32-7 FirstHelloWS]\$ pwd /kubews/FirstHelloWS [user1@ip-172-42-32-7 FirstHelloWS]\$ ls firsthellorc.yaml  firsthellosvc.yaml [user1@ip-172-42-32-7 FirstHelloWS]\$ █</pre> <pre> 1  kind: Service 2  apiVersion: v1 3  metadata: 4    name: firsthello-service 5  spec: 6    selector: 7      app: firsthello 8    ports: 9      - protocol: TCP 10        port: 8011 11        targetPort: 80 12        name: fh 13        type: NodePort </pre>
5	Deploy app onto Kubernetes cluster with create command as shown below

## Container orchestration with Kubernetes Hands-on Workbook

	<pre>[user1@ip-172-42-32-7 FirsthelloWS]\$ ls firsthellorc.yaml firsthellosvc.yaml [user1@ip-172-42-32-7 FirsthelloWS]\$ sudo kubectl create -f firsthellorc.yaml [sudo] password for user1: replicationcontroller/firsthello created [user1@ip-172-42-32-7 FirsthelloWS]\$ sudo kubectl get rc NAME      DESIRED   CURRENT   READY   AGE firsthello 1         1         1       13s [user1@ip-172-42-32-7 FirsthelloWS]\$ sudo kubectl get pods NAME          READY   STATUS    RESTARTS   AGE firsthello-f6md6 1/1     Running   0          17s mysql-6bfc69bb74-s6rf9 1/1     Running   0          5h [user1@ip-172-42-32-7 FirsthelloWS]\$</pre>
6	<p>Create service to expose deployed app using create command as shown below</p> <pre>[user1@ip-172-42-32-7 FirsthelloWS]\$ ls firsthellorc.yaml firsthellosvc.yaml [user1@ip-172-42-32-7 FirsthelloWS]\$ sudo kubectl create -f firsthellosvc.yaml service/firsthellosvc created [user1@ip-172-42-32-7 FirsthelloWS]\$ sudo kubectl get services NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE firsthellosvc  ClusterIP  10.99.124.224  &lt;none&gt;        80/TCP      10s kubernetes  ClusterIP  10.96.0.1    &lt;none&gt;        443/TCP     14d [user1@ip-172-42-32-7 FirsthelloWS]\$ sudo kubectl describe service/firsthellosvc Name:            firsthellosvc Namespace:       default Labels:          app=firsthello Annotations:    &lt;none&gt; Selector:        app=firsthello Type:           ClusterIP IP:             10.99.124.224 Port:           &lt;unset&gt;  80/TCP TargetPort:     80/TCP Endpoints:      172.17.0.4:80 Session Affinity: None Events:          &lt;none&gt; [user1@ip-172-42-32-7 FirsthelloWS]\$</pre> <p>Run command as shown below to get "firsthellosvc" service URL</p> <pre>[user1@ip-172-42-32-7 FirsthelloWS]\$ sudo /usr/local/bin/minikube service firsthellosvc --url http://172.42.32.7:31882 [user1@ip-172-42-32-7 FirsthelloWS]\$</pre> <p>You can see app UI as shown below</p>

## Container orchestration with Kubernetes Hands-on Workbook

You can check downloaded image and running container as shown below

```
[user@ip-172-42-32-7 FirsthelloWS]$ sudo docker images
REPOSITORY          TAG      IMAGE ID      CREATED       SIZE
msrinivascharan/docker-kubernetes   module1    b0340533957b  29 hours ago  132MB
nginx              latest    8b89e48b5f15  2 weeks ago   109MB
mysql              5.6      da785d6c6ed8  2 weeks ago   256MB
hello-world        latest    2cb0d9787c7d3  3 weeks ago   1.85kB
k8s.gcr.io/kube-proxy-amd64        v1.10.0   bfc21aadcc7d3  4 months ago  97MB
k8s.gcr.io/kube-controller-manager-amd64  v1.10.0   ad86dbed1555  4 months ago  148MB
k8s.gcr.io/kube-scheduler-amd64    v1.10.0   704ba848e69a  4 months ago  50.4MB
k8s.gcr.io/kube-apiserver-amd64   v1.10.0   af20925d51a3  4 months ago  225MB
k8s.gcr.io/etc-d amd64           3.1.12   52920ad46f5b  4 months ago  193MB
k8s.gcr.io/kube-addon-manager     v8.6     9c16409588eb  5 months ago  78.4MB
k8s.gcr.io/k8s-dns-dnsMasq-nanny-amd64  1.14.8   c2ce1fffb51ed  6 months ago  41MB
k8s.gcr.io/k8s-dns-sidecar-amd64   1.14.8   6f7f2dc7fab5  6 months ago  42.2MB
k8s.gcr.io/k8s-dns-kube-dns-amd64  1.14.8   80cc5ead4b547  6 months ago  50.5MB
k8s.gcr.io/pause-amd64            3.1     da86e6ba6cal  7 months ago  742kB
k8s.gcr.io/kubernetes-dashboard-amd64  v1.8.1   e94df2f21bc0c  7 months ago  121MB
gcr.io/k8s-minikube/storage-provisioner  v1.8.1   4689081edb10  8 months ago  80.8MB
wordpress             4.8-apache  fcf3e41b8864  8 months ago  408MB
jocatalin/kubernetes-bootcamp    v2      b6556396ebd4  24 months ago  211MB
gcr.io/google-samples/kubernetes-bootcamp  v1      8fafdbaf70e9  24 months ago  211MB
gcr.io/google-samples/gb-frontend      v4      e2b3e8542af7  2 years ago   512MB
gcr.io/google-samples/gb-redisslave   v1      5f026ddffa27  2 years ago   116MB
k8s.gcr.io/redis            e2e     e5e67996c442  3 years ago   419MB
[user@ip-172-42-32-7 FirsthelloWS]$
```

```
[user@ip-172-42-32-7 FirsthelloWS]$ sudo docker container ls
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              PORTS               NAMES
805db743fa8        c2ce1fffb51ed   "/dnsMasq-Nanny -v2."   About a minute ago   Up About a minute   0.0.0.0:44165->44165/tcp   k8s_dnsMasq_kube-dns_805db743fa8_default_0ae5cb43-8a66-11e8-9525-0a3f336079ae_962
38ce26f525af        msrinivascharan/docker-kubernetes   "python app.py"      7 minutes ago      Up 7 minutes      0.0.0.0:44166->44166/tcp   k8s_firsthello(firsthello-f6md6)_default_02ed0cca-9583-11e8-a645-0a3f336079ae_0
a0aa8e866172        k8s.gcr.io/pause-amd64:3.1   "/pause"           7 minutes ago      Up 7 minutes      0.0.0.0:44167->44167/tcp   k8s_POD(firsthello-f6md6)_0
c19d6fe7999        da785d6c6ed8   "docker-entrypoint.s..."  7 minutes ago      Up 7 minutes      0.0.0.0:44168->44168/tcp   k8s_POD(firsthello-f6md6)_1
a39364dc69ae        k8s.gcr.io/pause-amd64:3.1   "/pause"           5 hours ago       Up 5 hours       0.0.0.0:44169->44169/tcp   k8s_POD(firsthello-f6md6)_2
eba9237917c        4689081edb10   "/storage-provisioner"  5 hours ago       Up 5 hours       0.0.0.0:44170->44170/tcp   k8s_storage-provisioner(storage-provisioner)_0
feee7645f7af        bfc21aadcc7d3  "/usr/local/bin/kube..."  5 hours ago       Up 5 hours       0.0.0.0:44171->44171/tcp   k8s_kube-proxy(kube-proxy_5nf5r)_0
c85423a992a         k8s.gcr.io/pause-amd64:3.1   "/pause"           5 hours ago       Up 5 hours       0.0.0.0:44172->44172/tcp   k8s_POD(firsthello-f6md6)_3
8d2f9e0163a         k8s.gcr.io/pause-amd64:3.1   "/pause"           5 hours ago       Up 5 hours       0.0.0.0:44173->44173/tcp   k8s_POD(firsthello-f6md6)_4
653af0e9271d        k8s.gcr.io/pause-amd64:3.1   "/pause"           5 hours ago       Up 5 hours       0.0.0.0:44174->44174/tcp   k8s_POD(firsthello-f6md6)_5
df308f8bb5f         k8s.gcr.io/pause-amd64:3.1   "/pause"           5 hours ago       Up 5 hours       0.0.0.0:44175->44175/tcp   k8s_POD(firsthello-f6md6)_6
52e2983762d        ad86dbed1555   "kube-controller-man..."  5 hours ago       Up 5 hours       0.0.0.0:44176->44176/tcp   k8s_kube-controller-manager(minikube)_0
bc6ad6dab442        af20925d51a3   "kube-apiserver -ad..."  5 hours ago       Up 5 hours       0.0.0.0:44177->44177/tcp   k8s_kube-apiserver(minikube)_0
```

## Exercise 4.2 Deploying webserver on to Kubernetes cluster

## Scenario

Deploy Nginx on Kubernetes cluster

## Container orchestration with Kubernetes Hands-on Workbook

## To do tasks:

In this exercise, you will be:

- Deploy Nginx on Kubernetes cluster
- kubectl, clusters, the control plane, namespaces, pods, services, replication controllers, and labels.

**Note:** Alternatively, you can deploy Tomcat also using specifications available in /kubews/TomcatWS

## Steps

## Deploying Nginx

- 1 Create workspace “kubews” as shown below

```
[user1@ip-172-42-32-7 /]$ sudo mkdir kubews
[user1@ip-172-42-32-7 /]$ ls
bin  boot  data  dev  etc  home  kubews  lib  lib64  media  mnt  opt  proc  root  run  sbin  srv  sys  tmp  usr  var
[user1@ip-172-42-32-7 /]$ █
```

- 2

Add configuration file to create new namespace, command is “sudo vi development-ns.yaml”.

```
[user1@ip-172-42-32-7 kubews]$ pwd
/kubews
[user1@ip-172-42-32-7 kubews]$ sudo vi development-ns.yaml █
```

Add configurations as shown below

```
[user1@ip-172-42-32-7 kubews]$ sudo cat development-ns.yaml
kind: "Namespace"
apiVersion: "v1"
metadata:
  name: "development"
  labels:
    name: "development"
[user1@ip-172-42-32-7 kubews]$
```

Create namespace “Development” using command “sudo kubectl create -f development-ns.yaml” as shown below

	<pre>[user1@ip-172-42-32-7 kubews]\$ sudo kubectl create -f development-ns.yaml namespace/development created [user1@ip-172-42-32-7 kubews]\$ █</pre>
--	---

Check newly created namespace as shown below

	<pre>[user1@ip-172-42-32-7 kubews]\$ sudo kubectl get namespaces NAME      STATUS   AGE default   Active   1d development Active   2m kube-public Active   1d kube-system Active   1d [user1@ip-172-42-32-7 kubews]\$</pre>
--	---

3

Add configuration file to create new pod, command is “sudo vi nginx.yaml”.

	<pre>[user1@ip-172-42-32-7 kubews]\$ pwd /kubews [user1@ip-172-42-32-7 kubews]\$ sudo vi nginx.yaml</pre>
--	---

Add configurations as shown below

## Container orchestration with Kubernetes Hands-on Workbook

```
[user1@ip-172-42-32-7 kubews]$ sudo cat nginx.yaml
apiVersion: v1

kind: Pod

metadata:
  name: nginx

  labels:
    name: "development"

spec:
  containers:
    - name: nginx-server

      image: nginx

      ports:
        - containerPort: 80
[user1@ip-172-42-32-7 kubews]$ █
```

Create pod nginx using command “`sudo kubectl create -f nginx.yaml`” as shown below

```
[user1@ip-172-42-32-7 kubews]$ sudo kubectl create -f nginx.yaml
pod/nginx created
[user1@ip-172-42-32-7 kubews]$ █
```

Check newly created pod as shown below

```
[user1@ip-172-42-32-7 kubews]$ sudo kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
kubernetes-bootcamp-7799cbcb86-mjtxx   1/1     Running   0          3h
kubernetes-bootcamp-7799cbcb86-pgnfn   1/1     Running   0          3h
nginx         1/1     Running   0          1m
[user1@ip-172-42-32-7 kubews]$ █
```

## Container orchestration with Kubernetes Hands-on Workbook

```
[user1@ip-172-42-32-7 kubews]$ sudo kubectl describe pod/nginx
Name:           nginx
Namespace:      default
Node:          minikube/172.42.32.7
Start Time:    Thu, 19 Jul 2018 12:49:18 +0000
Labels:         name=development
Annotations:   <none>
Status:        Running
IP:            172.17.0.3
Containers:
  nginx-server:
    Container ID: docker://3819f183c0e1519b232eaf4e2f7a3b715c32edb607c5adbea26f31431bfa123a
    Image:          nginx
    Image ID:     docker-pullable://nginx@sha256:4a5573037f358b6cdaf2f3e8a9c33a5cf11bcd1675ca72ca76fbe5bd77d0d682
    Port:          80/TCP
    Host Port:    0/TCP
    State:        Running
      Started:   Thu, 19 Jul 2018 12:49:25 +0000
    Ready:        True
    Restart Count: 0
    Environment:  <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-d4bll (ro)
Conditions:
  Type      Status
  Initialized  True
  Ready      True
  PodScheduled  True
Volumes:
  default-token-d4bll:
    Type:      Secret (a volume populated by a Secret)
    SecretName: default-token-d4bll
    Optional:   false
  QoS Class:  BestEffort
  Node-Selectors: <none>
  Tolerations: node.kubernetes.io/not-ready:NoExecute for 300s
                node.kubernetes.io/unreachable:NoExecute for 300s
Events:
  Type  Reason          Age   From           Message
  ----  ----          ----  --           -----
  Normal Scheduled       4m   default-scheduler  Successfully assigned nginx to minikube
  Normal SuccessfulMountVolume 4m   kubelet, minikube  MountVolume.SetUp succeeded for volume "default-token-d4bll"
  Normal Pulling         4m   kubelet, minikube  pulling image "nginx"
  Normal Pulled          4m   kubelet, minikube  Successfully pulled image "nginx"
  Normal Created         4m   kubelet, minikube  Created container
  Normal Started         4m   kubelet, minikube  Started container
[user1@ip-172-42-32-7 kubews]$
```

Note: you can delete this pod using command “sudo kubectl delete pod nginx”. In the next step, we will create pod for Nginx using replication controller

4

Add configuration file to create new replication controller, command is “sudo vi nginxrc.yaml”.

```
[user1@ip-172-42-32-7 kubews]$ pwd
/kubews
[user1@ip-172-42-32-7 kubews]$ sudo vi nginxrc.yaml
[user1@ip-172-42-32-7 kubews]$
```

Add configurations as shown below

## Container orchestration with Kubernetes Hands-on Workbook

```
[user1@ip-172-42-32-7 kubews]$ sudo cat nginxrc.yaml
apiVersion: v1
kind: ReplicationController
metadata:
  name: rcnginx
  labels:
    name: "development"
spec:
  replicas: 1
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx
      ports:
        - containerPort: 80
```

Create replication controller nginx using command “`sudo kubectl create -f nginx.yaml`” as shown below

```
[user1@ip-172-42-32-7 kubews]$ sudo kubectl create -f nginxrc.yaml
replicationcontroller/rcnginx created
[user1@ip-172-42-32-7 kubews]$
```

Check newly created replication control “rcnginx” as shown below

## Container orchestration with Kubernetes Hands-on Workbook

```
[user1@ip-172-42-32-7 kubews]$ sudo kubectl get rc
NAME      DESIRED   CURRENT   READY     AGE
rcnginx   1         1         1         59s
[user1@ip-172-42-32-7 kubews]$ █
```

```
[user1@ip-172-42-32-7 kubews]$ sudo kubectl get pods
NAME                           READY   STATUS    RESTARTS   AGE
kubernetes-bootcamp-7799cbc86-mjtxx   1/1    Running   0          4h
kubernetes-bootcamp-7799cbc86-pgnfn   1/1    Running   0          4h
rcnginx-nlj7j                     1/1    Running   0          18m
[user1@ip-172-42-32-7 kubews]$ █
```

```
[user1@ip-172-42-32-7 kubews]$ sudo kubectl describe rc
Name:            rcnginx
Namespace:       default
Selector:        app=nginx
Labels:          name=development
Annotations:     <none>
Replicas:        1 current / 1 desired
Pods Status:    1 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:  app=nginx
  Containers:
    nginx:
      Image:      nginx
      Port:       80/TCP
      Host Port:  0/TCP
      Environment: <none>
      Mounts:     <none>
      Volumes:    <none>
  Events:
    Type     Reason     Age   From           Message
    ----     ----     --   --             -----
    Normal   SuccessfulCreate  3m   replication-controller  Created pod: rcnginx-nlj7j
[user1@ip-172-42-32-7 kubews]$ █
```

5

Scale out pods as shown below using command “`sudo kubectl scale --replicas=4 rc rcnginx`”

```
[user1@ip-172-42-32-7 kubews]$ sudo kubectl scale --replicas=4 rc rcnginx
replicationcontroller/rcnginx scaled
[user1@ip-172-42-32-7 kubews]$ sudo kubectl get pods
NAME                           READY   STATUS    RESTARTS   AGE
kubernetes-bootcamp-7799cbc86-mjtxx   1/1    Running   0          4h
kubernetes-bootcamp-7799cbc86-pgnfn   1/1    Running   0          4h
rcnginx-gmr4b                     0/1    ContainerCreating   0          3s
rcnginx-nlj7j                     1/1    Running   0          20m
rcnginx-pbvb                     0/1    ContainerCreating   0          3s
rcnginx-vrl58                     0/1    ContainerCreating   0          3s
[user1@ip-172-42-32-7 kubews]$ █
```

Little later, run get pods commands to see the updated status of pods

```
[user1@ip-172-42-32-7 kubews]$ sudo kubectl get pods
NAME                               READY   STATUS    RESTARTS   AGE
kubernetes-bootcamp-7799cbc86-mjtxx   1/1    Running   0          4h
kubernetes-bootcamp-7799cbc86-pgnfn   1/1    Running   0          4h
rcnginx-gmr4b                      1/1    Running   0          2m
rcnginx-hlj7j                      1/1    Running   0          22m
rcnginx-pbvbc                      1/1    Running   0          2m
rcnginx-vrl58                      1/1    Running   0          2m
[user1@ip-172-42-32-7 kubews]$
```

```
[user1@ip-172-42-32-7 kubews]$ sudo kubectl get rc
NAME      DESIRED   CURRENT   READY   AGE
rcnginx   4         4         4       25m
[user1@ip-172-42-32-7 kubews]$
```

6

Add configuration file to create new service “nginxsvc” , command is “sudo vi nginxsvc.yaml”.

```
[user1@ip-172-42-32-7 kubews]$ pwd
/kubews
[user1@ip-172-42-32-7 kubews]$ sudo vi nginxsvc.yaml
```

Add configurations as shown below

```
[user1@ip-172-42-32-7 kubews]$ sudo cat nginxsvc.yaml
apiVersion: v1

kind: Service

metadata:
  name: nginxsvc
  labels:
    app: nginx
spec:
  ports:
    - port: 80
      protocol: TCP
  selector:
    app: nginx
[user1@ip-172-42-32-7 kubews]$ █
```

Create service nginxsvc using command “`sudo kubectl create -f nginxsvc.yaml`” as shown below

```
[user1@ip-172-42-32-7 kubews]$ sudo kubectl create -f nginxsvc.yaml
service/nginxsvc created
```

Check newly created service as shown below

```
[user1@ip-172-42-32-7 kubews]$ sudo kubectl get service
NAME          TYPE        CLUSTER-IP   EXTERNAL-IP  PORT(S)       AGE
kubernetes    ClusterIP  10.96.0.1   <none>      443/TCP      1d
kubernetes-bootcamp  NodePort    10.99.186.215 <none>      8080:31216/TCP 8h
nginxsvc      ClusterIP  10.106.244.82  <none>      80/TCP      12s
[user1@ip-172-42-32-7 kubews]$ █
```

```
[user1@ip-172-42-32-7 kubews]$ sudo kubectl get service nginxsvc
NAME          TYPE        CLUSTER-IP   EXTERNAL-IP  PORT(S)       AGE
nginxsvc      ClusterIP  10.106.244.82  <none>      80/TCP      3m
[user1@ip-172-42-32-7 kubews]$ █
```

## Container orchestration with Kubernetes Hands-on Workbook

7	<p>Because every node in a Kubernetes cluster runs a kube-proxy, the kube-proxy watches the Kubernetes API server for the addition and removal of services. For each new service, kube-proxy opens a (randomly chosen) port on the local node. Any connections made to that port are proxied to one of the corresponding backend pods.</p> <p>Go to browser and check Nginx homepage as shown below</p>  <p>----- End of Exercise -----</p>

## Exercise 4.3 Deploying multi-tier app on to Kubernetes cluster

### Scenario

Deploy multi-tier web app on Kubernetes cluster

### To do tasks:

In this exercise, you will be:

- Build and deploy a simple, multi-tier web application using Kubernetes and Docker

This example consists of the following components:

1. A single-instance Redis master to store guestbook entries
2. Multiple replicated Redis instances to serve reads
3. Multiple web frontend instances

The guestbook application uses Redis to store its data. It writes its data to a Redis master instance and reads data from multiple Redis slave instances.

The guestbook application has a web frontend serving the HTTP requests written in PHP. It is configured to connect to the redis-master Service for write requests and the redis-slave service for Read requests.

### Steps

#### Deploying multi-tier web app

1	<p>Note: delete previous exercise related objects from cluster before starting this exercise</p> <p>Create workspace “/kubews/GuestbookAppWS” as shown below</p>
---	--

## Container orchestration with Kubernetes Hands-on Workbook

	<pre>[user1@ip-172-42-32-7 GuestbookAppWS]\$ pwd /kubews/GuestbookAppWS [user1@ip-172-42-32-7 GuestbookAppWS]\$ █</pre>
2	<p>Create configuration file “redis-master-deployment.yaml” as shown below, this is to specifies a Deployment controller that runs a single replica Redis master Pod</p> <pre>[user1@ip-172-42-32-7 GuestbookAppWS]\$ sudo cat redis-master-deployment.yaml apiVersion: apps/v1 # for versions before 1.9.0 use apps/v1beta2  kind: Deployment  metadata:    name: redis-master    labels:      app: redis  spec:    selector:      matchLabels:        app: redis        role: master        tier: backend    replicas: 1    template:      metadata:        labels:          app: redis          role: master          tier: backend      spec:        containers:        - name: master</pre>

## Container orchestration with Kubernetes Hands-on Workbook

	<pre> image: k8s.gcr.io/redis:e2e # or just image: redis  resources:  requests:  cpu: 100m  memory: 100Mi  ports:  - containerPort: 6379   </pre>
3	<p>Apply the Redis Master Deployment from file using command “<code>sudo kubectl create -f redis-master-deployment.yaml</code>”</p> <pre>[user1@ip-172-42-32-7 GuestbookAppWS]\$ sudo kubectl create -f redis-master-deployment.yaml deployment.apps/redis-master created [user1@ip-172-42-32-7 GuestbookAppWS]\$</pre> <p>You can check newly created pod as shown below.</p> <pre>[user1@ip-172-42-32-7 GuestbookAppWS]\$ sudo kubectl get pods NAME                               READY   STATUS    RESTARTS   AGE kubernetes-bootcamp-7799cbcb86-mjtxx   1/1    Running   1          1d kubernetes-bootcamp-7799cbcb86-pgnfn    1/1    Running   1          1d rcnginx-x-gmr4b                      1/1    Running   1          21h rcnginx-x-nlj7j                      1/1    Running   1          21h rcnginx-x-pbvbc                      1/1    Running   1          21h rcnginx-x-vrl58                      1/1    Running   1          21h redis-master-55db5f7567-npp5w         1/1    Running   0          5m [user1@ip-172-42-32-7 GuestbookAppWS]\$</pre> <p>Run the following command to view the logs from the Redis Master Pod</p> <pre>[user1@ip-172-42-32-7 GuestbookAppWS]\$ sudo kubectl logs -f redis-master-55db5f7567-npp5w [1] 20 Jul 10:45:24.924 # Server started, Redis version 2.8.19 [1] 20 Jul 10:45:24.924 # WARNING you have Transparent Huge Pages (THP) support enabled in your kernel. This will create latency and memory usage issues with Redis. To fix this issue run the command 'echo never &gt; /sys/kernel/mm/transparent_hugepage/enabled' as root, and add it to your /etc/rc.local in order to retain the setting after a reboot. Redis must be restarted after THP is disabled. [1] 20 Jul 10:45:24.924 # WARNING: The TCP backlog setting of 511 cannot be enforced because /proc/sys/net/core/somaxconn is set to the lower value 128. [1] 20 Jul 10:45:24.924 * The server is now ready to accept connections on port 6379</pre>

4

Create configuration file “redis-master-service.yaml” as shown below, this is to apply a Service to proxy the traffic to the Redis master Pod.

```
[user1@ip-172-42-32-7 GuestbookAppWS]$ sudo cat redis-master-service.yaml
apiVersion: v1

kind: Service

metadata:

  name: redis-master

  labels:

    app: redis

    role: master

    tier: backend

spec:

  ports:

    - port: 6379

      targetPort: 6379

  selector:

    app: redis

    role: master

    tier: backend
```

-

Apply the Redis Master Service from the above redis-master-service.yaml file, command is “`sudo kubectl create -f redis-master-service.yaml`”

```
[user1@ip-172-42-32-7 GuestbookAppWS]$ sudo kubectl create -f redis-master-service.yaml
service/redis-master created
[user1@ip-172-42-32-7 GuestbookAppWS]$
```

Check newly created services as shown below

## Container orchestration with Kubernetes Hands-on Workbook

	<pre>[user1@ip-172-42-32-7 GuestbookAppWS]\$ sudo kubectl get services NAME           TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE kubernetes     ClusterIP  10.96.0.1   &lt;none&gt;        443/TCP         2d kubernetes-bootcamp NodePort   10.99.186.215 &lt;none&gt;        8080:31216/TCP  1d nginxsvc       ClusterIP  10.106.244.82 &lt;none&gt;        80/TCP          22h redis-master    ClusterIP  10.109.178.54 &lt;none&gt;        6379/TCP         1m [user1@ip-172-42-32-7 GuestbookAppWS]\$</pre> <p>You can run describe command on newly created service</p> <p><b>Note:</b> This manifest file creates a Service named redis-master with a set of labels that match the labels previously defined, so the Service routes network traffic to the Redis master Pod.</p>
5	<p>Create configuration file “redis-slave-deployment.yaml” as shown below, this is to create Redis slaves</p> <pre>[user1@ip-172-42-32-7 GuestbookAppWS]\$ sudo cat redis-slave-deployment.yaml apiVersion: apps/v1 # for versions before 1.9.0 use apps/v1beta2  kind: Deployment  metadata:    name: redis-slave    labels:      app: redis  spec:    selector:      matchLabels:        app: redis        role: slave        tier: backend    replicas: 2    template:      metadata:        labels:</pre>

```
labels:  
  app: redis  
  role: slave  
  tier: backend  
  
spec:  
  containers:  
    - name: slave  
      image: gcr.io/google_samples/gb-redisslave:v1  
  resources:  
    requests:  
      cpu: 100m  
      memory: 100Mi  
  env:  
    - name: GET_HOSTS_FROM  
      value: env  
  ports:  
    - containerPort: 6379
```

Create the Redis Slave Deployment from the file, command is “`sudo kubectl create -f redis-slave-deployment.yaml`”

```
[user1@ip-172-42-32-7 GuestbookAppWS]$ sudo kubectl create -f redis-slave-deployment.yaml  
deployment.apps/redis-slave created  
[user1@ip-172-42-32-7 GuestbookAppWS]$
```

Check newly created Redis slave pods as shown below

## Container orchestration with Kubernetes Hands-on Workbook

	<pre>[user1@ip-172-42-32-7 GuestbookAppWS]\$ sudo kubectl get pods NAME                               READY   STATUS    RESTARTS   AGE kubernetes-bootcamp-7799cbcb86-mjtxx  1/1     Running   1          1d kubernetes-bootcamp-7799cbcb86-pgnfn  1/1     Running   1          1d rcnginx-gmr4b                      1/1     Running   1          22h rcnginx-nlj7j                      1/1     Running   1          23h rcnginx-pbvbc                      1/1     Running   1          22h rcnginx-vrl58                       1/1     Running   1          22h redis-master-55db5f7567-npp5w        1/1     Running   0          1h redis-slave-584c66c5b5-9vkmz        1/1     Running   0          2m redis-slave-584c66c5b5-j69v5        1/1     Running   0          2m [user1@ip-172-42-32-7 GuestbookAppWS]\$</pre>
6	Create configuration file “redis-slave-service.yaml” as shown below, this is to create Redis slave service.

## Container orchestration with Kubernetes Hands-on Workbook

```
[user1@ip-172-42-32-7 GuestbookAppWS]$ sudo cat redis-slave-service.yaml
apiVersion: v1

kind: Service

metadata:
  name: redis-slave
  labels:
    app: redis
    role: slave
    tier: backend
spec:
  ports:
  - port: 6379
  selector:
    app: redis
    role: slave
    tier: backend
[user1@ip-172-42-32-7 GuestbookAppWS]$ █
```

Create the Redis Slave Service from the above redis-slave-service.yaml file command is “`sudo kubectl create -f redis-slave-service.yaml`”

```
[user1@ip-172-42-32-7 GuestbookAppWS]$ sudo kubectl create -f redis-slave-service.yaml
service/redis-slave created
[user1@ip-172-42-32-7 GuestbookAppWS]$
```

Check the newly created service as shown below

```
[user1@ip-172-42-32-7 GuestbookAppWS]$ sudo kubectl get service
NAME           TYPE      CLUSTER-IP   EXTERNAL-IP  PORT(S)        AGE
kubernetes     ClusterIP  10.96.0.1   <none>       443/TCP       2d
kubernetes-bootcamp  NodePort   10.99.186.215 <none>       8080:31216/TCP 1d
nginxsvc       ClusterIP  10.106.244.82 <none>       80/TCP        22h
redis-master   ClusterIP  10.109.178.54 <none>       6379/TCP      38m
redis-slave    ClusterIP  10.111.101.87 <none>       6379/TCP      2m
[user1@ip-172-42-32-7 GuestbookAppWS]$ █
```

7	<p>Create configuration file “frontend-deployment.yaml” as shown below, this is to create Guest book app frontend.</p> <pre>[user1@ip-172-42-32-7 GuestbookAppWS]\$ sudo cat frontend-deployment.yaml apiVersion: apps/v1 # for versions before 1.9.0 use apps/v1beta2  kind: Deployment  metadata:   name: frontend    labels:     app: guestbook  spec:   selector:     matchLabels:       app: guestbook       tier: frontend   replicas: 3   template:     metadata:       labels:         app: guestbook</pre>

```
tier: frontend

spec:

  containers:
    - name: php-redis
      image: gcr.io/google-samples/gb-frontend:v4
      resources:
        requests:
          cpu: 100m
          memory: 100Mi
      env:
        - name: GET_HOSTS_FROM
          value: env
      ports:
        - containerPort: 80
[user1@ip-172-42-32-7 GuestbookAppWS]$ █
```

Apply the frontend Deployment from the frontend-deployment.yaml file, command is “`sudo kubectl create -f frontend-deployment.yaml`”

```
[user1@ip-172-42-32-7 GuestbookAppWS]$ sudo kubectl create -f frontend-deployment.yaml
deployment.apps/frontend created
[user1@ip-172-42-32-7 GuestbookAppWS]$ █
```

Check newly created pods as shown below

## Container orchestration with Kubernetes Hands-on Workbook

	<pre>[user1@ip-172-42-32-7 GuestbookAppWS]\$ sudo kubectl get pods NAME                               READY   STATUS    RESTARTS   AGE frontend-6d797fc7dd-2q27s          1/1    Running   0          2m frontend-6d797fc7dd-6r2vw          1/1    Running   0          2m frontend-6d797fc7dd-ph5ln          1/1    Running   0          2m kubernetes-bootcamp-7799c9cb86-mjtxx 1/1    Running   1          1d kubernetes-bootcamp-7799c9cb86-pgnfn 1/1    Running   1          1d rcnginx-gmr4b                      1/1    Running   1          23h rcnginx-nlj7j                      1/1    Running   1          23h rcnginx-pbvb                       1/1    Running   1          23h rcnginx-vrl58                       1/1    Running   1          23h redis-master-55db5f7567-npp5w       1/1    Running   0          2h redis-slave-584c66c5b5-9vkmz        1/1    Running   0          40m redis-slave-584c66c5b5-j69v5        1/1    Running   0          40m [user1@ip-172-42-32-7 GuestbookAppWS]\$ █</pre> <pre>[user1@ip-172-42-32-7 GuestbookAppWS]\$ sudo kubectl get pods -l app=guestbook -l tier=frontend NAME                               READY   STATUS    RESTARTS   AGE frontend-6d797fc7dd-2q27s          1/1    Running   0          4m frontend-6d797fc7dd-6r2vw          1/1    Running   0          4m frontend-6d797fc7dd-ph5ln          1/1    Running   0          4m [user1@ip-172-42-32-7 GuestbookAppWS]\$ █</pre>
8	<p>Note: The redis-slave and redis-master Services you applied are only accessible within the container cluster because the default type for a Service is ClusterIP. ClusterIP provides a single IP address for the set of Pods the Service is pointing to. This IP address is accessible only within the cluster.</p> <p>Note: Minikube can only expose Services through NodePort.</p> <p>Create configuration file “frontend-service.yaml” as shown below, this is to create Guest book app frontend service .</p>

## Container orchestration with Kubernetes Hands-on Workbook

```
[user1@ip-172-42-32-7 GuestbookAppWS]$ sudo cat frontend-service.yaml
apiVersion: v1

kind: Service

metadata:
  name: frontend
  labels:
    app: guestbook
    tier: frontend

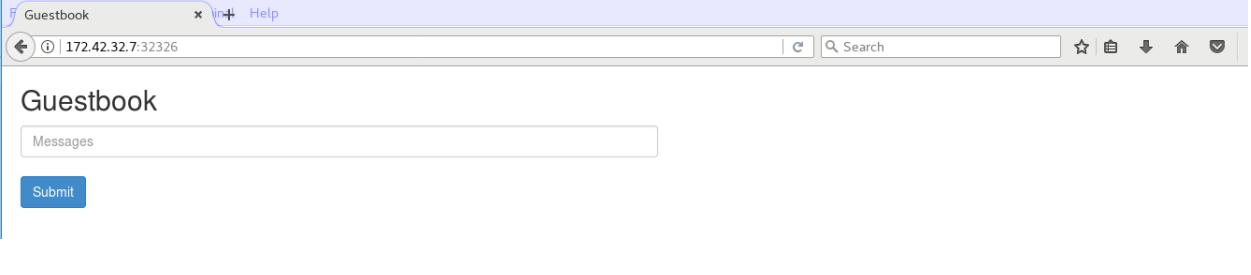
spec:
  # comment or delete the following line if you want to use a LoadBalancer
  type: NodePort
  # if your cluster supports it, uncomment the following to automatically create
  # an external load-balanced IP for the frontend service.
  # type: LoadBalancer
  ports:
    - port: 80
  selector:
    app: guestbook
    tier: frontend
```

Apply the frontend Service from the frontend-service.yaml file, command is “`sudo kubectl create -f frontend-service.yaml`”

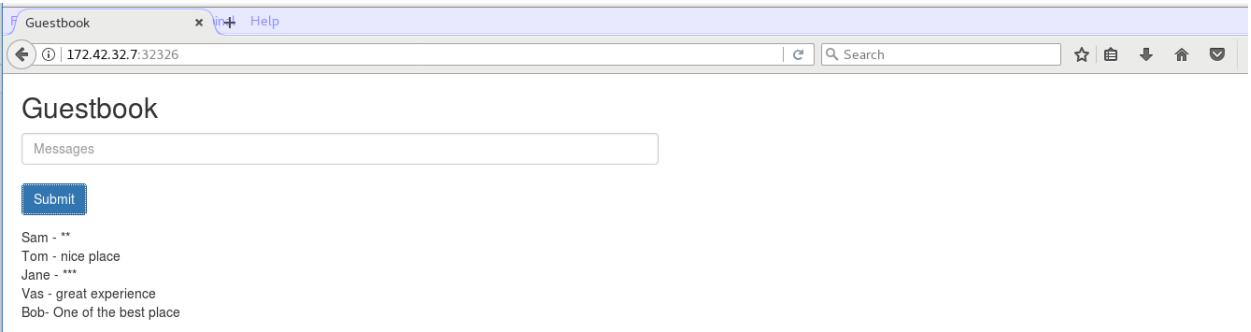
```
[user1@ip-172-42-32-7 GuestbookAppWS]$ sudo kubectl create -f frontend-service.yaml
service/frontend created
[user1@ip-172-42-32-7 GuestbookAppWS]$
```

Check newly created front end service as shown below

## Container orchestration with Kubernetes Hands-on Workbook

	<pre>[user1@ip-172-42-32-7 GuestbookAppWS]\$ sudo kubectl get service NAME           TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE frontend       NodePort   10.97.192.180 &lt;none&gt;        8018:30840/TCP  2m kubernetes     ClusterIP  10.96.0.1    &lt;none&gt;        443/TCP         2d kubernetes-bootcamp NodePort   10.99.186.215 &lt;none&gt;        8080:31216/TCP  1d nginxsvc       ClusterIP  10.106.244.82  &lt;none&gt;        80/TCP          23h redis-master   ClusterIP  10.109.178.54 &lt;none&gt;        6379/TCP        1h redis-slave    ClusterIP  10.111.101.87 &lt;none&gt;        6379/TCP        41m [user1@ip-172-42-32-7 GuestbookAppWS]\$</pre>
	<p>Suppose your guestbook app has been running for a while, and it gets a sudden burst of publicity. You decide it would be a good idea to add more web servers to your frontend. You can do this easily since your servers are defined as a service that uses a Deployment controller.</p>
9	<p>Check App status in cluster as shown below</p> <pre>[user1@ip-172-42-32-7 GuestbookAppWS]\$ sudo kubectl get po,svc -o wide NAME             READY   STATUS    RESTARTS   AGE   IP           NODE pod/frontend-6dd96bb6f7-9stb4  1/1    Running   0          5m   172.17.0.9  minikube pod/frontend-6dd96bb6f7-ktfxx  1/1    Running   0          5m   172.17.0.7  minikube pod/frontend-6dd96bb6f7-sm4fv  1/1    Running   0          5m   172.17.0.8  minikube pod/redis-master-55db5f7567-2j9gp 1/1    Running   1          29m  172.17.0.4  minikube pod/redis-slave-6894998d77-4knrb 1/1    Running   0          12m  172.17.0.5  minikube pod/redis-slave-6894998d77-wmpjt 1/1    Running   0          12m  172.17.0.6  minikube  NAME           TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE   SELECTOR service/frontend   NodePort   10.108.169.43 &lt;none&gt;        80:32326/TCP  1m   app=guestbook,tier=frontend service/kubernetes ClusterIP  10.96.0.1    &lt;none&gt;        443/TCP         5d   &lt;none&gt; service/redis-master ClusterIP  10.107.157.196 &lt;none&gt;        6379/TCP        28m  app=redis,role=master,tier=backend service/redis-slave  ClusterIP  10.108.45.211  &lt;none&gt;        6379/TCP        8m   app=redis,role=slave,tier=backend [user1@ip-172-42-32-7 GuestbookAppWS]\$</pre>
10	<p>Run command “sudo /usr/local/bin/minikube service frontend --url” to get App frontend service URL</p> <pre>[user1@ip-172-42-32-7 GuestbookAppWS]\$ sudo /usr/local/bin/minikube service frontend --url http://172.42.32.7:32326 [user1@ip-172-42-32-7 GuestbookAppWS]\$</pre>
12	<p>Enter URL in the web browser, access Guest Book App home page</p> 

## Container orchestration with Kubernetes Hands-on Workbook

	<p>Enter messages and submit as shown below</p> 
13	<p>We can use <code>kubectl exec</code> to get a shell to a running Container</p> <p>Get the Redis master pod name, run command “<code>sudo kubectl exec -ti redis-master-55db5f7567-2j9gp bash</code>” as shown below (running) to query Redis guest book database.</p> <pre>[user@ip-172-42-32-7 GuestbookAppWS]\$ sudo kubectl exec -ti redis-master-55db5f7567-2j9gp bash [ root@redis-master-55db5f7567-2j9gp:/data ]\$ date Mon Jul 23 13:04:59 UTC 2018 [ root@redis-master-55db5f7567-2j9gp:/data ]\$</pre> <p>Query database as shown below</p> <pre>[ root@redis-master-55db5f7567-2j9gp:/data ]\$ redis-cli KEYS "*" 1) "messages" [ root@redis-master-55db5f7567-2j9gp:/data ]\$ redis-cli KEYS "messages" 1) "messages" [ root@redis-master-55db5f7567-2j9gp:/data ]\$ redis-cli MGET "messages" 1) ",Sam - **,Tom - nice place,Jane - ***,Vas - great experience,Bob- One of the best place" [ root@redis-master-55db5f7567-2j9gp:/data ]\$</pre>
14	<p>Suppose your guestbook app has been running for a while, and it gets a sudden burst of publicity. You decide it would be a good idea to add more web servers to your frontend. You can do this easily since your servers are defined as a service that uses a Deployment controller.</p> <p>Run the command “<code>kubectl scale deployment frontend --replicas=5</code>” to scale up the number of frontend Pods</p>

## Container orchestration with Kubernetes Hands-on Workbook

	<pre>[user1@ip-172-42-32-7 GuestbookAppWS]\$ sudo kubectl get pods NAME READY STATUS RESTARTS AGE frontend-6dd96bb6f7-9stb4 1/1 Running 0 35m frontend-6dd96bb6f7-ktfxx 1/1 Running 0 35m frontend-6dd96bb6f7-sm4fv 1/1 Running 0 35m redis-master-55db5f7567-2j9gp 1/1 Running 1 59m redis-slave-6894998d77-4knrb 1/1 Running 0 42m redis-slave-6894998d77-wnpjt 1/1 Running 0 42m [user1@ip-172-42-32-7 GuestbookAppWS]\$ sudo kubectl scale deployment frontend --replicas=5 deployment.extensions/frontend scaled [user1@ip-172-42-32-7 GuestbookAppWS]\$ sudo kubectl get pods NAME READY STATUS RESTARTS AGE frontend-6dd96bb6f7-87lwv 1/1 Running 0 3s frontend-6dd96bb6f7-9stb4 1/1 Running 0 35m frontend-6dd96bb6f7-ktfxx 1/1 Running 0 35m frontend-6dd96bb6f7-s8r5b 1/1 Running 0 3s frontend-6dd96bb6f7-sm4fv 1/1 Running 0 35m redis-master-55db5f7567-2j9gp 1/1 Running 1 59m redis-slave-6894998d77-4knrb 1/1 Running 0 42m redis-slave-6894998d77-wnpjt 1/1 Running 0 42m [user1@ip-172-42-32-7 GuestbookAppWS]\$</pre>
15	<p>Run the command “<code>kubectl scale deployment frontend --replicas=2</code>” to scale down the number of frontend Pods</p> <pre>[user1@ip-172-42-32-7 GuestbookAppWS]\$ sudo kubectl get pods NAME READY STATUS RESTARTS AGE frontend-6dd96bb6f7-87lwv 1/1 Running 0 2m frontend-6dd96bb6f7-9stb4 1/1 Running 0 38m frontend-6dd96bb6f7-ktfxx 1/1 Running 0 38m frontend-6dd96bb6f7-s8r5b 1/1 Running 0 2m frontend-6dd96bb6f7-sm4fv 1/1 Running 0 38m redis-master-55db5f7567-2j9gp 1/1 Running 1 1h redis-slave-6894998d77-4knrb 1/1 Running 0 45m redis-slave-6894998d77-wnpjt 1/1 Running 0 45m [user1@ip-172-42-32-7 GuestbookAppWS]\$ sudo kubectl scale deployment frontend --replicas=2 deployment.extensions/frontend scaled [user1@ip-172-42-32-7 GuestbookAppWS]\$ sudo kubectl get pods NAME READY STATUS RESTARTS AGE frontend-6dd96bb6f7-9stb4 1/1 Running 0 38m frontend-6dd96bb6f7-s8r5b 0/1 Terminating 0 2m frontend-6dd96bb6f7-sm4fv 1/1 Running 0 38m redis-master-55db5f7567-2j9gp 1/1 Running 1 1h redis-slave-6894998d77-4knrb 1/1 Running 0 45m redis-slave-6894998d77-wnpjt 1/1 Running 0 45m [user1@ip-172-42-32-7 GuestbookAppWS]\$ sudo kubectl get pods NAME READY STATUS RESTARTS AGE frontend-6dd96bb6f7-9stb4 1/1 Running 0 38m frontend-6dd96bb6f7-sm4fv 1/1 Running 0 38m redis-master-55db5f7567-2j9gp 1/1 Running 1 1h redis-slave-6894998d77-4knrb 1/1 Running 0 45m redis-slave-6894998d77-wnpjt 1/1 Running 0 45m [user1@ip-172-42-32-7 GuestbookAppWS]\$</pre>
16	<p>You can use below mentioned commands to remove Guest book app from cluster</p>

- |  |   |
|--|---|
|  | <ol style="list-style-type: none"><li>1. kubectl delete deployment -l app=redis</li><li>2. kubectl delete service -l app=redis</li><li>3. kubectl delete deployment -l app=guestbook</li><li>4. kubectl delete service -l app=guestbook</li></ol> |
|--|---|

----- End of Exercise -----

## Exercise 4.4 Deploying MySQL onto Kubernetes cluster with persistent volumes (optional)

### Scenario

Deploy MySQL database to Kubernetes cluster

### To do tasks:

In this exercise, you will be:

- Create Persistent Volume Claims and Persistent Volumes
- Create a Secret
- Deploy MySQL

### Steps

#### Deploy MySQL database Kubernetes single node cluster

1	Create workspace as shown below  [user1@ip-172-42-32-7 MySQLws]\$ pwd /kubews/MySQLws [user1@ip-172-42-32-7 MySQLws]\$ █
2	Create and update “mysql-pv.yaml” configuration file as shown below, to create persistent volume and claims

	<pre>kind: PersistentVolume apiVersion: v1 metadata:   name: mysql-pv-volume   labels:     type: local spec:   storageClassName: manual   capacity:     storage: 8Gi   accessModes:     - ReadWriteOnce   hostPath:     path: "/mnt/data" --- apiVersion: v1 kind: PersistentVolumeClaim metadata:   name: mysql-pv-claim spec:   storageClassName: manual   accessModes:     - ReadWriteOnce   resources:     requests:       storage: 8Gi</pre>
3	Create and update “mysql-deployment.yaml” configuration file as shown below, to deploy MySQL

## Container orchestration with Kubernetes Hands-on Workbook

	<pre> apiVersion: v1 kind: Service metadata:   name: mysql spec:   ports:     - port: 3306   selector:     app: mysql   clusterIP: None --- apiVersion: apps/v1 # for versions before 1.9.0 use apps/v1beta2 kind: Deployment metadata:   name: mysql spec:   selector:     matchLabels:       app: mysql   strategy:     type: Recreate   template:     metadata:       labels:         app: mysql     spec:       containers:         - image: mysql:5.6           name: mysql           env:             # Use secret in real usage             - name: MYSQL_ROOT_PASSWORD               valueFrom:                 secretKeyRef:                   name: mysql-pass                   key: password           ports:             - containerPort: 3306               name: mysql           volumeMounts:             - name: mysql-persistent-storage               mountPath: /var/lib/mysql       volumes:         - name: mysql-persistent-storage           persistentVolumeClaim:             claimName: mysql-pv-claim </pre>
4	<p>Create secret to store MySQL password. Command is “<code>sudo kubectl create secret generic mysql-pass --from-literal=password=password</code>”</p> <pre>[user1@ip-172-42-32-7 MySQLws]\$ sudo kubectl create secret generic mysql-pass --from-literal=password=password secret/mysql-pass created [user1@ip-172-42-32-7 MySQLws]\$</pre>
5	<p>Create PV and PVC using command “<code>sudo kubectl create -f mysql-pv.yaml</code>”</p>

## Container orchestration with Kubernetes Hands-on Workbook

	<pre>[user1@ip-172-42-32-7 MySQLws]\$ sudo kubectl create -f mysql-pv.yaml persistentvolume/mysql-pv-volume created persistentvolumeclaim/mysql-pv-claim created [user1@ip-172-42-32-7 MySQLws]\$</pre>
6	<p>Create MySQL pod and service using command “<code>sudo kubectl create -f mysql-deployment.yaml</code>”</p> <pre>[user1@ip-172-42-32-7 MySQLws]\$ sudo kubectl create -f mysql-deployment.yaml service/mysql created deployment.apps/mysql created [user1@ip-172-42-32-7 MySQLws]\$</pre>
7	<p>Get Pod, service, PV and PVC and secrets details as shown below, command is “<code>sudo kubectl get po,svc,pv,pvc,secrets -o wide</code>”</p> <pre>[user1@ip-172-42-32-7 MySQLws]\$ sudo kubectl get po,svc,pv,pvc,secrets -o wide NAME           READY   STATUS    RESTARTS   AGE   IP          NODE pod/mysql-6bfc69bb74-qbpq4  1/1    Running   0        2m   172.17.0.4   minikube  NAME          TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE   SELECTOR service/kubernetes ClusterIP  10.96.0.1   &lt;none&gt;      443/TCP   9d   &lt;none&gt; service/mysql   ClusterIP  None        &lt;none&gt;      3306/TCP  2m   app=mysql  NAME          CAPACITY   ACCESS MODES  RECLAIM POLICY   STATUS   CLAIM           STORAGECLASS   REASON   AGE persistentvolume/mysql-pv-volume  8Gi        RWO          Retain       Bound   default/mysql-pv-claim  manual        3m  NAME          STATUS     VOLUME      CAPACITY   ACCESS MODES  STORAGECLASS   AGE persistentvolumeclaim/mysql-pv-claim  Bound   mysql-pv-volume  8Gi        RWO          manual        3m  NAME          TYPE        DATA        AGE secret/default-token-d4bll  kubernetes.io/service-account-token  3        9d secret/mysql-pass        Opaque      1          5m [user1@ip-172-42-32-7 MySQLws]\$</pre>
8	<p>We can use <code>kubectl exec</code> to get a shell to a running Container</p> <p>Get the MySQL pod name, run command “<code>sudo kubectl exec -ti mysql-6bfc69bb74-qbpq4 bash</code>” as shown below (running) to access MySQL database.</p> <pre>[user1@ip-172-42-32-7 MySQLws]\$ sudo kubectl get pods NAME           READY   STATUS    RESTARTS   AGE mysql-6bfc69bb74-qbpq4  1/1    Running   0        4m [user1@ip-172-42-32-7 MySQLws]\$ sudo kubectl exec -ti mysql-6bfc69bb74-qbpq4 bash root@mysql-6bfc69bb74-qbpq4:/#</pre>
9	<p>After step 8, run below mentioned SQL commands to create database and store some data</p> <ol style="list-style-type: none"> <li>1. <code>mysql -u root -p</code> (password is “password”)</li> <li>2. <code>show databases;</code></li> </ol>

## Container orchestration with Kubernetes Hands-on Workbook

	<p>3. CREATE DATABASE pets;</p> <p>4. USE pets;</p> <p>5. CREATE TABLE pet (name VARCHAR(20), owner VARCHAR(20), species VARCHAR(20), sex CHAR(1), birth DATE, death DATE);</p> <p>6. show tables;</p> <p>7. describe pet;</p>
10	<p>Insert data into newly created pet table using below mentioned SQL commands</p> <pre>INSERT INTO pet VALUES ('Puffball','Diane','hamster','f','1999-03-30',NULL); INSERT INTO pet VALUES ('Ulal','Diane','Fox','m','1999-03-30',NULL); INSERT INTO pet VALUES ('Lovely','Diane','Parrot','f','1999-03-30',NULL); INSERT INTO pet VALUES ('Jiss','Diane','Lab','f','1999-03-30',NULL); INSERT INTO pet VALUES ('James','Diane','Beagle','f','1999-03-30',NULL); INSERT INTO pet VALUES ('Snowey','Diane','Husky','f','1999-03-30',NULL); INSERT INTO pet VALUES ('Short','Diane','Spaniel','f','1999-03-30',NULL); INSERT INTO pet VALUES ('Chief','Diane','Sheppard','f','1999-03-30',NULL);</pre> <p>Check table data using command “select * from pet”</p>
11	<p>Come out of MySQL shell using exit command, delete pods and other related Objects using below mentioned commands.</p> <ol style="list-style-type: none"> <li>1. kubectl delete deployment,svc mysql</li> <li>2. kubectl delete pvc mysql-pv-claim</li> <li>3. kubectl delete pv mysql-pv-volume</li> </ol>
12	<p>Repeat steps from 5 to 8 and check database we created in step 10 and 11 existed or not in newly deployed MySQL instance. Because we used persistent volumes, database exists.</p>

----- End of Exercise -----

## Exercise x.x Setup Kubernetes cluster using Kubeadm

### Scenario

Kubernetes cluster setup

### To do tasks:

In this exercise

- Kubernetes cluster setup using Kubeadm

### Steps

Kubernetes cluster setup	
1	On master node:

## Container orchestration with Kubernetes Hands-on Workbook

	<p>Install Docker using below steps</p> <pre> 1] Go to https://download.docker.com/linux/centos/7/x86_64/stable/Packages/, download 2] cd /home/user1/Downloads 3] sudo yum install docker-ce-18.06.1.ce-3.el7.x86_64.rpm 4] cd .. 5] sudo systemctl start docker 6] sudo docker run hello-world 7] sudo docker images </pre>
2	<p>On master node:</p> <p>Step 1: cd to directory yum.repos.d using command “cd /etc/yum.repos.d”</p> <p>Step 2: Create file kubernetes.repo using subl and add content to that file as shown below</p>  <pre> [kubernetes] name=Kubernetes baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-x86_64 enabled=1 gpgcheck=1 gpgkey=https://packages.cloud.google.com/yum/doc/yum-key.gpg https://packages.cloud.google.com/yum/doc/rpm-package-key.gpg exclude=kube* </pre> <p>Step 3: Install Kubelet, kubectl, Kubeadm with command “yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes”</p> <p>Step 4: Run command “sudo systemctl enable kubelet” to enable kubelet</p> <p>Step 5: Run command “sudo systemctl start kubelet” to start kubelet</p>
3	<p>On master node:</p> <p>Step 1: Run command “sudo subl /proc/sys/net/bridge/bridge-nf-call-iptables” to override 0 with 1</p> <p>Step 2: Run command “sudo kubeadm init” to initialize master node</p> <p>Step 3: Copy Kubeadm token to join nodes with master</p> <p>“kubeadm join 172.42.1.55:6443 --token pe8nhr.5yavfa6dufosbrc7 --discovery-token-ca-cert-hash sha256:d1d55b2e83f7cf101de8b82e2258c843cbf0c9a5c8ab7307119835fb9b77ff08”</p> <p>Step 4: Run commands as shown below to use cluster</p> <pre> 1] mkdir -p \$HOME/.kube 2] sudo cp -i /etc/kubernetes/admin.conf \$HOME/.kube/config 3] sudo chown \$(id -u):\$(id -g) \$HOME/.kube/config </pre>
4	<p>On master node:</p>

## Container orchestration with Kubernetes Hands-on Workbook

---

	<p>Step 1: cd to /home/user1/kubews/ClustersetupWS</p> <p>Step 2: Create new yaml file, command is “sudo subl podntwrkadon.yaml”</p> <p>Step 3: Add configurations to newly created yaml file from <a href="https://raw.githubusercontent.com/coreos/flannel/v0.10.0/Documentation/kube-flannel.yml">https://raw.githubusercontent.com/coreos/flannel/v0.10.0/Documentation/kube-flannel.yml</a></p> <p>Step 4:</p>
--	---